

Deep learning for timeseries

by Sanjeev Gupta

Deep learning for timeseries

by Sanjeev Gupta

✓ A temperature-forecasting example

Problem:

Given a timeseries of hourly measurements of various atmospheric parameters, **predict** the temperature 24 hours in the future.

Goal:

To demonstrate that RNNs perform better than Dense or Conv Networks for time-series data

```
!wget https://s3.amazonaws.com/keras-datasets/jena_climate_2009_2016.csv.zip
!unzip jena_climate_2009_2016.csv.zip
```

```
--2022-06-03 16:33:17-- https://s3.amazonaws.com/keras-datasets/jena_climate_2009_2016.csv.zip
Resolving s3.amazonaws.com (s3.amazonaws.com)... 52.217.203.80
Connecting to s3.amazonaws.com (s3.amazonaws.com)|52.217.203.80|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 13565642 (13M) [application/zip]
Saving to: 'jena_climate_2009_2016.csv.zip'

jena_climate_2009_2 100%[=====] 12.94M 57.6MB/s in 0.2s

2022-06-03 16:33:18 (57.6 MB/s) - 'jena_climate_2009_2016.csv.zip' saved [13565642/13565642]

Archive: jena_climate_2009_2016.csv.zip
  inflating: jena_climate_2009_2016.csv
  inflating: __MACOSX/._jena_climate_2009_2016.csv
```

```
import tensorflow as tf
import os
import pandas as pd
```

```
zip_path = tf.keras.utils.get_file(
    origin='https://storage.googleapis.com/tensorflow/tf-keras-datasets/jena_climate_2009_2016.csv.zip',
    fname='jena_climate_2009_2016.csv.zip',
    extract=True)
csv_path, _ = os.path.splitext(zip_path)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/jena_climate_2009_2016.csv.zip
13574144/13568290 [=====] - 0s 0us/step
13582336/13568290 [=====] - 0s 0us/step
```

Inspecting the data of the Jena weather dataset

```
df = pd.read_csv(csv_path)
# Slice [start:stop:step], starting from index 5 take every 6th record.
#df = df[5::6]

date_time = pd.to_datetime(df.pop('Date Time'), format='%d.%m.%Y %H:%M:%S')

df.head()
```

```


p      T      Tpot      Tdew      rh      VPmax      VPact      VPdef      sh      H2OC      rho      ww      max. ww      wd
(mbar) (degC) (K) (degC) (%) (mbar) (mbar) (mbar) (g/kg) (mmol/mol) (g/m**3) (m/s) (m/s) (deg)

0  996.52 -8.02  265.40 -8.90  93.3   3.33   3.11   0.22   1.94   3.12  1307.75  1.03  1.75  152.3
1  996.57 -8.41  265.01 -9.28  93.4   3.23   3.02   0.21   1.89   3.03  1309.80  0.72  1.50  136.1
2  996.53 -8.51  264.91 -9.31  93.9   3.21   3.01   0.20   1.88   3.02  1310.24  0.19  0.63  171.6
3  996.51 -8.31  265.12 -9.07  94.2   3.26   3.07   0.19   1.92   3.08  1309.19  0.34  0.50  198.0
4  996.51 -8.27  265.15 -9.04  94.1   3.27   3.08   0.19   1.92   3.09  1309.00  0.32  0.63  214.3
```

Double-click (or enter) to edit

```
# **Visualize the data**
plot_cols = ['T (degC)', 'p (mbar)', 'rho (g/m**3)']
plot_features = df[plot_cols]
plot_features.index = date_time
_ = plot_features.plot(subplots=True)

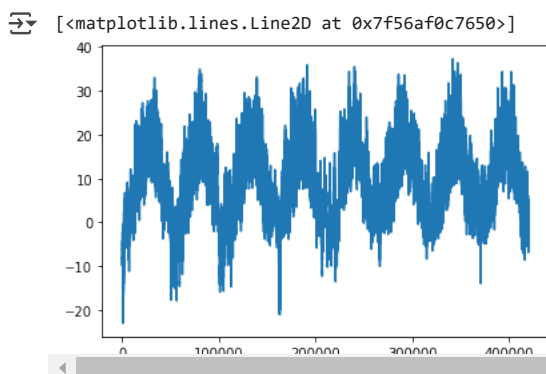
# Lets zoom in to a smaller time-window
plot_features = df[plot_cols][:480]
plot_features.index = date_time[:480]
_ = plot_features.plot(subplots=True)
```



Let's just focus on the temperature and plot.

Note: Data is recorded every 10 minutes, you get $24 \times 6 = 144$ data points per day

```
# Plotting the temperature timeseries
from matplotlib import pyplot as plt
temperature = df['T (degC)'] # extract out the temperature data
plt.plot(range(len(temperature)), temperature)
```



```
# Convert data to np array
import numpy as np
raw_data = np.array(df) #raw_data is a np array
print(raw_data.shape)
# Q: how many features are there? A:
```


```
(420551, 14)
```

Double-click (or enter) to edit

```
# Plotting the first 10 days of the temperature timeseries
plt.plot(range(1440), temperature[:1440])
print(temperature[:5])
```

```
0 -8.02
1 -8.41
2 -8.51
3 -8.31
4 -8.27
```

Name: T (degC), dtype: float64



```
# Compute the number of samples we'll use for each data split
num_train_samples = int(0.5 * len(raw_data)) #Q: What is len(raw_data) #A:
num_val_samples = int(0.25 * len(raw_data))
num_test_samples = len(raw_data) - num_train_samples - num_val_samples #Q: apx how many test samples? #A: apx 25%
print("num_train_samples:", num_train_samples)
print("num_val_samples:", num_val_samples)
print("num_test_samples:", num_test_samples)
```

```
num_train_samples: 210275
num_val_samples: 105137
num_test_samples: 105139
```

✓ Preparing the data

Normalizing the data

```
# Important: Let's not repeat mistakes
mean = raw_data[:num_train_samples].mean(axis=0) # Q: axis=0 means? Is mean a vector? #A:
std = raw_data[:num_train_samples].std(axis=0)
```


temporal

```
raw_data -= mean
raw_data /= std
```

```
# Important Q: On what data was the mean and std computed? #A:
# Q: On what data is standarization tranformation applied? #A:
```

```
# plot the normalized temperature data
print(raw_data[:,1]) # temperature data
plt.plot(raw_data[:1440,1])
#Q: What is the difference between the below plot and the prev one? #A: Scale, visible diff in y-axis
#Q: Does the plot look like it has 0 mean? Yes/No/Why? #A:
```

```
[-1.92080466 -1.96527448 -1.976677 ... -1.36664229 -1.48864923
-1.55592409]
[<matplotlib.lines.Line2D at 0x7f56b22d5cd0>]
```



We have time series data from sensors/web/...

But how do we get datasets? What is the input? What is the label?

To make apt datasets, We will use the a keras [utility](#) called

```
timeseries_dataset_from_array()
```

So first, let's see how it works

(Q: Do you know of any other utility for some other kind of dataset?)

```
import numpy as np
from tensorflow import keras
int_sequence = np.arange(20)
print(f"Original Timeseries sequence: {int_sequence}\n")
print(f"data = int_sequence[:-3]={int_sequence[:-3]}\n")
print(f"targets= int_sequence[3:] = {int_sequence[3:]}\n")

# returns sequences and corresponding targets
dummy_dataset = keras.utils.timeseries_dataset_from_array(
    data=int_sequence[:-3],
    targets=int_sequence[3:],
    # The next 3 arguments are for you to manipulate the data (input)
    sequence_length=4,
    sampling_rate = 3,
    sequence_stride = 2, # stride applied to both data and targets
    # batch_size=2, # see in next cell
)

for inputs, targets in dummy_dataset:
    # print(inputs, targets)
    print(inputs.shape)
    for i in range(inputs.shape[0]): # range(batch_size)
        print([int(x) for x in inputs[i]], int(targets[i]))
```

Original Timeseries sequence: [0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19]

data = int_sequence[:-3]=[0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16]

targets= int_sequence[3:] = [3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19]

(3, 4)

[0, 3, 6, 9] 3

[2, 5, 8, 11] 5

[4, 7, 10, 13] 7

Q: We want the following:

[0, 1, 2] : 3

[1, 2, 3] : 4

[2, 3, 4] : 5

...

...

Start coding or [generate](#) with AI.

Q: We want the following:

[0, 2, 4] : 3

[1, 3, 5] : 4

[2, 4, 6] : 5

...

...

Q: We want the following:

[0, 3, 6, 9] : 3

[2, 5, 8, 11] : 5

[4, 7, 10, 12] : 7

...

...

Q: What if the target list is too small?

```
int_sequence = np.arange(20)
print(f"Original Timeseries sequence: {int_sequence}")
print(f"data = int_sequence[:-3]={int_sequence[:-3]}")
print(f"targets= int_sequence[3:] = {int_sequence[3:]}")

dummy_dataset = keras.utils.timeseries_dataset_from_array(
    data=int_sequence[:-3],      #Q: What is the effect of the 3 #A: maintains delay of 3 between la
    targets=int_sequence[3:],
    # The next 3 arguments are for you to manipulate the data (input)
    sequence_length=5,
    # sampling_rate = 2,
    # sequence_stride = 3,
    # batch_size=2,
    # shuffle = True
)

for inputs, targets in dummy_dataset:
    print(inputs)
    print(targets)
    print(f"\n inputs.shape = {inputs.shape}")
    for i in range(inputs.shape[0]):
        print([int(x) for x in inputs[i]], int(targets[i]))
```

```
Original Timeseries sequence: [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19]
data = int_sequence[:-3]=[ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16]
targets= int_sequence[3:] = [ 3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19]
tf.Tensor(
[[ 0  1  2  3  4]
 [ 1  2  3  4  5]
 [ 2  3  4  5  6]
 [ 3  4  5  6  7]
 [ 4  5  6  7  8]
 [ 5  6  7  8  9]
 [ 6  7  8  9 10]
 [ 7  8  9 10 11]
 [ 8  9 10 11 12]
 [ 9 10 11 12 13]
[10 11 12 13 14]
[11 12 13 14 15]
[12 13 14 15 16]], shape=(13, 5), dtype=int64)
tf.Tensor([ 3  4  5  6  7  8  9 10 11 12 13 14 15], shape=(13,), dtype=int64)

inputs.shape = (13, 5)
[0, 1, 2, 3, 4] 3
[1, 2, 3, 4, 5] 4
[2, 3, 4, 5, 6] 5
[3, 4, 5, 6, 7] 6
[4, 5, 6, 7, 8] 7
[5, 6, 7, 8, 9] 8
[6, 7, 8, 9, 10] 9
[7, 8, 9, 10, 11] 10
[8, 9, 10, 11, 12] 11
[9, 10, 11, 12, 13] 12
[10, 11, 12, 13, 14] 13
[11, 12, 13, 14, 15] 14
[12, 13, 14, 15, 16] 15
```

Now that we understand how the utility works, let's use it to make our dataset

Remember: Data is recorded every 10 minutes, you get $24 \times 6 = 144$ data points per day

Instantiating datasets for training, validation, and testing

Observations will be sampled at one data point per hour: we will only keep one data point out of 6.
sampling_rate = 6 # every 6th sample of the original sequence is a sample after an hour

sequence_length = 120 #Q: How many days worth data in a sequence? #A: x Days

The target for a sequence will be the temperature 24 hours after the end of the sequenc
delay = sampling_rate * (sequence_length + 24 - 1)

batch_size = 256

```

train_dataset = keras.utils.timeseries_dataset_from_array(
    raw_data[:-delay],          # Q: What is this argument? #A: data
    targets=temperature[delay:],
    sampling_rate=sampling_rate,
    sequence_length=sequence_length,
    shuffle=True,
    batch_size=batch_size,
    start_index=0,
    end_index=num_train_samples)

val_dataset = keras.utils.timeseries_dataset_from_array(
    raw_data[:-delay],
    targets=temperature[delay:],
    sampling_rate=sampling_rate,
    sequence_length=sequence_length,
    shuffle=True,
    batch_size=batch_size,
    start_index=num_train_samples,
    end_index=num_train_samples + num_val_samples)

test_dataset = keras.utils.timeseries_dataset_from_array(
    raw_data[:-delay],
    targets=temperature[delay:],
    sampling_rate=sampling_rate,
    sequence_length=sequence_length,
    shuffle=True,
    batch_size=batch_size,
    start_index=num_train_samples + num_val_samples)

```

Q: Do we have to normalise the test data or is it already done?

```

# Inspecting the output of one of our datasets
for samples, targets in train_dataset:
    print("samples shape:", samples.shape)
    print("targets shape:", targets.shape)
    break

```

```

↗ samples shape: (256, 120, 14)
  targets shape: (256,)

```

✓ A common-sense, non-machine-learning baseline

Computing the common-sense baseline MAE

Prediction based on the assumption that the temperature changes periodically.

Hence, the temperature prediction of 24 hrs later should be equal to the temperature right now

```

def evaluate_naive_method(dataset):
    total_abs_err = 0.
    samples_seen = 0
    for samples, targets in dataset:
        #Q: What is the shape of samples? What is idx 1? # A: (256, 120, 14), ; idx 1 means temperature. See dataframe.
        preds = samples[:, -1, 1] * std[1] + mean[1] #Q: Why mult by std and add with mean?
        #Q: What are the next 2 lines doing?
        total_abs_err += np.sum(np.abs(preds - targets))
        samples_seen += samples.shape[0]
    return total_abs_err / samples_seen

print(f"Validation MAE: {evaluate_naive_method(val_dataset):.2f}")
print(f"Test MAE: {evaluate_naive_method(test_dataset):.2f}")

```

```

↗ Validation MAE: 2.44
  Test MAE: 2.62

```

The Test MAE is 2.62. It means our naive model is expected to give a prediction which is off by 2.62 deg on average.

✓ Let's try a dense NN

We want to see how a densely connected network would perform on the problem.

```

# **Training and evaluating a densely connected model**
from tensorflow import keras
from tensorflow.keras import layers

```

```

from tensorflow.keras.utils import plot_model

# define input layer: relate to how you would pass image data to a dense layer
# Remember: Shape of an input batch: (256, 128, 14)
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1])) #Q: What is raw_data.shape[-1]? #A: Last index, i.e., no. of features

x = layers.Flatten()(inputs) #Q: Why flatten? # A: Input is in the form of matrix
x = layers.Dense(16, activation="relu")(x)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)

callbacks = [
    keras.callbacks.ModelCheckpoint("jena_dense.keras",
                                    save_best_only=True)
]
model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
history = model.fit(train_dataset,
                    epochs=10,
                    validation_data=val_dataset,
                    callbacks=callbacks)

model = keras.models.load_model("jena_dense.keras")
print(f"Test MAE: {model.evaluate(test_dataset)[1]:.2f}")

```

```

Epoch 1/10
819/819 [=====] - 43s 49ms/step - loss: 15.3177 - mae: 3.0231 - val_loss: 12.9447 - val_mae: 2.8549
Epoch 2/10
819/819 [=====] - 40s 49ms/step - loss: 10.1642 - mae: 2.5066 - val_loss: 10.9849 - val_mae: 2.6413
Epoch 3/10
819/819 [=====] - 41s 50ms/step - loss: 9.3269 - mae: 2.4004 - val_loss: 12.2271 - val_mae: 2.7855
Epoch 4/10
819/819 [=====] - 41s 49ms/step - loss: 8.8321 - mae: 2.3372 - val_loss: 10.7274 - val_mae: 2.6012
Epoch 5/10
819/819 [=====] - 40s 48ms/step - loss: 8.3972 - mae: 2.2775 - val_loss: 11.8348 - val_mae: 2.7474
Epoch 6/10
819/819 [=====] - 40s 49ms/step - loss: 8.1313 - mae: 2.2416 - val_loss: 10.6627 - val_mae: 2.5916
Epoch 7/10
819/819 [=====] - 41s 50ms/step - loss: 7.8942 - mae: 2.2082 - val_loss: 11.0262 - val_mae: 2.6427
Epoch 8/10
819/819 [=====] - 40s 49ms/step - loss: 7.6486 - mae: 2.1763 - val_loss: 11.0131 - val_mae: 2.6403
Epoch 9/10
819/819 [=====] - 40s 49ms/step - loss: 7.4516 - mae: 2.1465 - val_loss: 11.2592 - val_mae: 2.6802
Epoch 10/10
819/819 [=====] - 40s 49ms/step - loss: 7.3133 - mae: 2.1287 - val_loss: 11.5753 - val_mae: 2.7012
405/405 [=====] - 14s 33ms/step - loss: 2765.6377 - mae: 6.2533
Test MAE: 6.25


```

Plotting results

```

import matplotlib.pyplot as plt
loss = history.history["mae"]
val_loss = history.history["val_mae"]
epochs = range(1, len(loss) + 1)
plt.figure()
plt.plot(epochs, loss, "bo", label="Training MAE")
plt.plot(epochs, val_loss, "b", label="Validation MAE")
plt.title("Training and validation MAE")
plt.legend()
plt.show()

```



The dense network has not done very well. Achieves a val loss of 2.6.

Why?

- The hypothesis space is inappropriate
 - perhaps too large.... needle in a haystack

- looking for a window-wise global pattern
- An example of the powerful nature of
 - good feature engineering
 - domain knowledge

✓ Let's try a 1D convolutional model


Now lets try a CNN.

But we don't have an image?

- We don't necessarily need one

We will use 1D convolutions.

- try to learn local patterns in the signal/sequence



[Nice gif](#)

```
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.utils import plot_model
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))

x = layers.Conv1D(8, 24, activation="relu", input_shape=(120,14))(inputs)
x = layers.MaxPooling1D(2)(x)
x = layers.Conv1D(8, 12, activation="relu")(x)
x = layers.MaxPooling1D(2)(x)
x = layers.Conv1D(8, 6, activation="relu")(x)
x = layers.GlobalAveragePooling1D()(x)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)
callbacks = [
    keras.callbacks.ModelCheckpoint("jena_conv.keras",
                                    save_best_only=True)
]
model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
plot_model(model, show_shapes=True)
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-2-64d4cf0dc158> in <module>
      2 from tensorflow.keras import layers
      3 from tensorflow.keras.utils import plot_model
----> 4 inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
      5
      6 x = layers.Conv1D(8, 24, activation="relu", input_shape=(120,14))(inputs)

NameError: name 'sequence_length' is not defined
```

```
history = model.fit(train_dataset,
                    epochs=10,
                    validation_data=val_dataset,
                    callbacks=callbacks)
```

```
model = keras.models.load_model("jena_conv.keras")
print(f"Test MAE: {model.evaluate(test_dataset)[1]:.2f}")
```

```
Epoch 1/10
819/819 [=====] - 43s 51ms/step - loss: 21.6386 - mae: 3.6381 - val_loss: 16.0439 - val_mae: 3.1527
Epoch 2/10
819/819 [=====] - 42s 51ms/step - loss: 15.2469 - mae: 3.1052 - val_loss: 14.3108 - val_mae: 2.9796
Epoch 3/10
819/819 [=====] - 42s 51ms/step - loss: 13.8960 - mae: 2.9554 - val_loss: 13.9624 - val_mae: 2.9398
Epoch 4/10
```



```

819/819 [=====] - 43s 52ms/step - loss: 12.9435 - mae: 2.8534 - val_loss: 14.6220 - val_mae: 2.9818
Epoch 5/10
819/819 [=====] - 43s 52ms/step - loss: 12.2739 - mae: 2.7790 - val_loss: 13.5374 - val_mae: 2.8815
Epoch 6/10
819/819 [=====] - 43s 52ms/step - loss: 11.7057 - mae: 2.7121 - val_loss: 13.9518 - val_mae: 2.9199
Epoch 7/10
819/819 [=====] - 43s 52ms/step - loss: 11.1627 - mae: 2.6484 - val_loss: 13.2862 - val_mae: 2.8556
Epoch 8/10
819/819 [=====] - 42s 51ms/step - loss: 10.7152 - mae: 2.5972 - val_loss: 14.8524 - val_mae: 3.0112
Epoch 9/10
819/819 [=====] - 42s 51ms/step - loss: 10.4060 - mae: 2.5580 - val_loss: 14.5681 - val_mae: 2.9907
Epoch 10/10
819/819 [=====] - 42s 51ms/step - loss: 10.1176 - mae: 2.5244 - val_loss: 15.4485 - val_mae: 3.0636
405/405 [=====] - 14s 33ms/step - loss: 21641.4277 - mae: 12.9118
Test MAE: 12.91

```

✓ A first recurrent baseline

Built-in RNN layers: a simple example

There are three built-in RNN layers in Keras:

- **keras.layers.SimpleRNN**, a fully-connected RNN where the output from previous timestep is to be fed to next timestep.
- **keras.layers.GRU**, first proposed in Cho et al., 2014.
- **keras.layers.LSTMs**, first proposed in Hochreiter & Schmidhuber, 1997.

Let's see how an RNN performs.

We will have a detailed look at the architecture shortly.

```

inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.LSTM(16)(inputs)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)

callbacks = [
    keras.callbacks.ModelCheckpoint("jena_lstm.keras",
                                    save_best_only=True)
]
model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
history = model.fit(train_dataset,
                    epochs=10,
                    validation_data=val_dataset,
                    callbacks=callbacks)

model = keras.models.load_model("jena_lstm.keras")
print(f"Test MAE: {model.evaluate(test_dataset)[1]:.2f}")

```

```

↩ Epoch 1/10
819/819 [=====] - 127s 149ms/step - loss: 36.1634 - mae: 4.3527 - val_loss: 11.3400 - val_mae: 2.5657
Epoch 2/10
819/819 [=====] - 122s 148ms/step - loss: 10.5254 - mae: 2.5235 - val_loss: 9.5128 - val_mae: 2.4038
Epoch 3/10
819/819 [=====] - 122s 149ms/step - loss: 9.7208 - mae: 2.4238 - val_loss: 9.5953 - val_mae: 2.4128
Epoch 4/10
819/819 [=====] - 122s 149ms/step - loss: 9.3853 - mae: 2.3779 - val_loss: 9.7360 - val_mae: 2.4320
Epoch 5/10
819/819 [=====] - 123s 149ms/step - loss: 9.1134 - mae: 2.3424 - val_loss: 9.5317 - val_mae: 2.4168
Epoch 6/10
819/819 [=====] - 122s 148ms/step - loss: 8.9050 - mae: 2.3153 - val_loss: 9.5940 - val_mae: 2.4203
Epoch 7/10
819/819 [=====] - 121s 147ms/step - loss: 8.7368 - mae: 2.2925 - val_loss: 9.8174 - val_mae: 2.4390
Epoch 8/10
819/819 [=====] - 120s 146ms/step - loss: 8.6120 - mae: 2.2743 - val_loss: 9.7383 - val_mae: 2.4391
Epoch 9/10
819/819 [=====] - 122s 149ms/step - loss: 8.4287 - mae: 2.2489 - val_loss: 9.8159 - val_mae: 2.4562

```

The test MAE is 2.57

- Slightly better than its Dense counterpart!
- Easy to use. Just a line of code.
- We can do much better. But we will not spend much time on training RNN based architecture. Instead we will invest that time on Transformers.

✓ Understanding recurrent neural networks

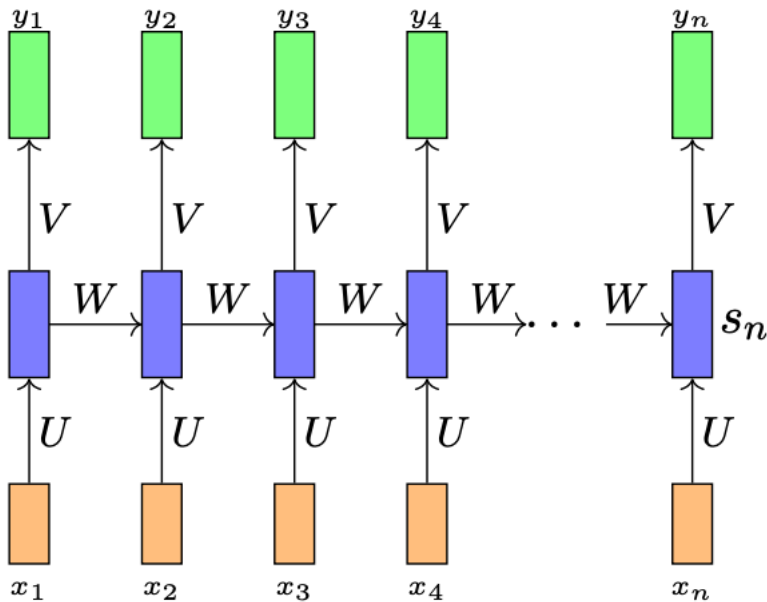
Sequence Learning problems:

- Inputs are no longer independent
- Input sizes not fixed. Ex- Length of sentences

So an RNN should take care of:

- Dependence between inputs
- Variable no. of inputs
- Function executed at each time must be the same

NumPy implementation of a simple RNN




$$s_i = \sigma(Ux_i + Ws_{i-1} + b)$$

$$y_i = \mathcal{O}(Vs_i + c)$$

or

$$y_i = f(x_i, s_{i-1}, W, U, V, b, c)$$



```
import numpy as np
timesteps = 100  #:number of timesteps in the input sequence
input_features = 32
```

```

output_features = 64
inputs = np.random.random((timesteps, input_features))
state_t = np.zeros((output_features,))
U = np.random.random((output_features, input_features))
W = np.random.random((output_features, output_features))
b = np.random.random((output_features,))
successive_outputs = []
for input_t in inputs:
    output_t = np.tanh(np.dot(U, input_t) + np.dot(W, state_t) + b)
    successive_outputs.append(output_t)
    state_t = output_t
final_output_sequence = np.stack(successive_outputs, axis=0) # Q: seq/vec-to-seq/vec #A: seq-to-seq due to stack

```

Notice the sequential nature of computation. We will discuss this later.

✓ A recurrent layer in Keras

An RNN layer that can process sequences of any length

```


num_features = 14
inputs = keras.Input(shape=(None, num_features))
outputs = layers.SimpleRNN(16)(inputs)
model = keras.Model(inputs, outputs)
print(outputs.shape)
model.summary()
plot_model(model, show_shapes=True)

```

↗ (None, 16)
Model: "model_2"

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, None, 14)]	0
simple_rnn_1 (SimpleRNN)	(None, 16)	496

=====
Total params: 496
Trainable params: 496
Non-trainable params: 0



$16 \times 14 + 16 \times 16 + 16$

An RNN layer that returns only its last output step

```

num_features = 14
steps = 120
inputs = keras.Input(shape=(steps, num_features))
outputs = layers.SimpleRNN(16, return_sequences=False)(inputs) # return_seq = False
model = keras.Model(inputs, outputs)
print(outputs.shape) # see output
model.summary()
plot_model(model, show_shapes=True)

```

(None, 16)
Model: "model_4"

Layer (type)	Output Shape	Param #
input_6 (InputLayer)	[(None, 120, 14)]	0
simple_rnn_2 (SimpleRNN)	(None, 16)	496

=====
Total params: 496
Trainable params: 496
Non-trainable params: 0

input_6	input:	[(None, 120, 14)]	[(None, 120, 14)]
InputLayer	output:		

simple_rnn_2	input:	(None, 120, 14)	(None, 16)
SimpleRNN	output:		

```
num_features = 14
steps = 120
inputs = keras.Input(shape=(steps, num_features))
outputs = layers.LSTM(16, return_sequences=False)(inputs) # return_seq = False
model = keras.Model(inputs, outputs)
print(outputs.shape) # see output
model.summary()
plot_model(model, show_shapes=True)
```

(None, 16)
Model: "model_3"

Layer (type)	Output Shape	Param #
input_5 (InputLayer)	[(None, 120, 14)]	0
lstm_1 (LSTM)	(None, 16)	1984

=====
Total params: 1,984
Trainable params: 1,984
Non-trainable params: 0

input_5	input:	[(None, 120, 14)]	[(None, 120, 14)]
InputLayer	output:		

lstm_1	input:	(None, 120, 14)	(None, 16)
LSTM	output:		

An RNN layer that returns its full output sequence

```
num_features = 14
steps = 120
inputs = keras.Input(shape=(steps, num_features))
outputs = layers.SimpleRNN(16, return_sequences=True)(inputs)
print(outputs.shape)
model = keras.Model(inputs, outputs)
model.summary()
plot_model(model, show_shapes=True)
```

(None, 120, 16)

Stacking RNN layers




Figure 15-7. Deep RNN (left) unrolled through time (right)

```
inputs = keras.Input(shape=(steps, num_features))
x = layers.SimpleRNN(16, return_sequences=True)(inputs) # Q: can return_seq be False here? # A: NO. It will throw error.
x = layers.SimpleRNN(16, return_sequences=True)(x)
outputs = layers.SimpleRNN(16)(x)
model = keras.Model(inputs, outputs)
print(outputs.shape)
model.summary()
```

(None, 16)
Model: "model_9"

Layer (type)	Output Shape	Param #
input_11 (InputLayer)	[(None, 120, 14)]	0
simple_rnn_5 (SimpleRNN)	(None, 120, 16)	496
simple_rnn_6 (SimpleRNN)	(None, 120, 16)	528
simple_rnn_7 (SimpleRNN)	(None, 16)	528

=====
Total params: 1,552
Trainable params: 1,552
Non-trainable params: 0
=====

16*16 + 16*16 + 16

✓ Advanced use of recurrent neural networks

- Recurrent dropout
- Stacking recurrent layers
- Bi-directional RNNs

✓ Using recurrent dropout to fight overfitting

Key Idea: The dropout pattern should remain the same across time.

```
# Training and evaluating a dropout-regularized LSTM
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.LSTM(32, recurrent_dropout=0.25)(inputs)
```

```

x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)

callbacks = [
    keras.callbacks.ModelCheckpoint("jena_lstm_dropout.keras",
                                    save_best_only=True)
]
model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
history = model.fit(train_dataset,
                    epochs=50,
                    validation_data=val_dataset,
                    callbacks=callbacks)


```

⚠ WARNING:tensorflow:Layer lstm_1 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as a fallback.

Epoch 1/50
819/819 [=====] - 334s 405ms/step - loss: 27.6691 - mae: 3.8794 - val_loss: 9.8712 - val_mae: 2.4367
Epoch 2/50
819/819 [=====] - 333s 406ms/step - loss: 14.7891 - mae: 2.9858 - val_loss: 9.2836 - val_mae: 2.3735
Epoch 3/50
819/819 [=====] - 334s 408ms/step - loss: 14.0565 - mae: 2.9057 - val_loss: 9.0759 - val_mae: 2.3511
Epoch 4/50
819/819 [=====] - 335s 409ms/step - loss: 13.3965 - mae: 2.8380 - val_loss: 8.9423 - val_mae: 2.3318
Epoch 5/50
819/819 [=====] - 335s 408ms/step - loss: 12.9484 - mae: 2.7880 - val_loss: 8.9122 - val_mae: 2.3240
Epoch 6/50
819/819 [=====] - 337s 411ms/step - loss: 12.4860 - mae: 2.7364 - val_loss: 9.2643 - val_mae: 2.3582
Epoch 7/50
819/819 [=====] - 334s 408ms/step - loss: 12.1242 - mae: 2.7004 - val_loss: 9.3036 - val_mae: 2.3678
Epoch 8/50
819/819 [=====] - 334s 407ms/step - loss: 11.7401 - mae: 2.6545 - val_loss: 9.2084 - val_mae: 2.3540
Epoch 9/50
819/819 [=====] - 335s 409ms/step - loss: 11.5305 - mae: 2.6305 - val_loss: 9.2848 - val_mae: 2.3767
Epoch 10/50
819/819 [=====] - 337s 411ms/step - loss: 11.3035 - mae: 2.6045 - val_loss: 9.4993 - val_mae: 2.3928
Epoch 11/50
819/819 [=====] - 336s 410ms/step - loss: 11.1175 - mae: 2.5820 - val_loss: 9.3366 - val_mae: 2.3672
Epoch 12/50
819/819 [=====] - 333s 406ms/step - loss: 10.8707 - mae: 2.5565 - val_loss: 9.4777 - val_mae: 2.3874
Epoch 13/50
819/819 [=====] - 333s 407ms/step - loss: 10.7547 - mae: 2.5408 - val_loss: 9.9690 - val_mae: 2.4502
Epoch 14/50
819/819 [=====] - 334s 407ms/step - loss: 10.6095 - mae: 2.5264 - val_loss: 9.8857 - val_mae: 2.4407
Epoch 15/50
819/819 [=====] - 335s 409ms/step - loss: 10.5469 - mae: 2.5158 - val_loss: 9.9061 - val_mae: 2.4366
Epoch 16/50
819/819 [=====] - 337s 412ms/step - loss: 10.3363 - mae: 2.4920 - val_loss: 9.9430 - val_mae: 2.4488
Epoch 17/50
483/819 [=====>.....] - ETA: 2:10 - loss: 10.2551 - mae: 2.4851

How is the loss defined?

Especially for a seq-to-seq RNN model?




✓ Stacking recurrent layers

We've already seen this. But repeating an example to show GRU.

```
# Training and evaluating a dropout-regularized, stacked GRU model
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.GRU(32, recurrent_dropout=0.5, return_sequences=True)(inputs) #Q: why ret_seq =True? #A: Stacked GRUs
x = layers.GRU(32, recurrent_dropout=0.5)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)

callbacks = [
    keras.callbacks.ModelCheckpoint("jena_stacked_gru_dropout.keras",
                                    save_best_only=True)
]
model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
# history = model.fit(train_dataset,
#                     epochs=50,
#                     validation_data=val_dataset,
#                     callbacks=callbacks)
# model = keras.models.load_model("jena_stacked_gru_dropout.keras")
# print(f"Test MAE: {model.evaluate(test_dataset)[1]:.2f}")
```

✓ Using bidirectional RNNs



Key Idea: Learn temporal patterns in both directions

Found to work well with text data.

```
# Training and evaluating a bidirectional LSTM
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.Bidirectional(layers.LSTM(16))(inputs)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)

model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
# history = model.fit(train_dataset,
#                     epochs=10,
#                     validation_data=val_dataset)
model = keras.Model(inputs, outputs)
print(outputs.shape)
model.summary()
```

(None, 1)
Model: "model_10"

Layer (type)	Output Shape	Param #
=====		
input_9 (InputLayer)	[(None, 120, 14)]	0
bidirectional_1 (Bidirectional)	(None, 32)	3968
dense_4 (Dense)	(None, 1)	33

```

=====
Total params: 4,001
Trainable params: 4,001
Non-trainable params: 0
=====

```

Let's compare it with a single LSTM layer

```

# Training and evaluating a bidirectional LSTM
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
# x = layers.Bidirectional(layers.LSTM(16))(inputs)
x = layers.LSTM(16)(inputs)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)

model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
# history = model.fit(train_dataset,
#                     epochs=10,
#                     validation_data=val_dataset)
model = keras.Model(inputs, outputs)
print(outputs.shape)
model.summary()

```

```

↻ (None, 1)
Model: "model_12"

```

Layer (type)	Output Shape	Param #
input_10 (InputLayer)	[(None, 120, 14)]	0
lstm_5 (LSTM)	(None, 16)	1984
dense_5 (Dense)	(None, 1)	17

```

=====
Total params: 2,001
Trainable params: 2,001
Non-trainable params: 0
=====

```

Note on runtime performance of RNNs:

- Small RNNs run faster on CPU
- Large RNNs can benefit from GPU
- GPU support not offered by cuDNN when using functionalities that are not optimized. E.g. Recurrent Dropout.
- Alternative: Unroll = True
- Caveat: Sequence length must be known apriori

```

inputs = keras.Input(shape=(sequence_length, num_features))
x = layers.LSTM(32, recurrent_dropout=0.2, unroll=True)(inputs)

```

✓ Assignment 4

Given a sequence of atmospheric measurements, predict a forecast of the **Density** after **2** days.

Solve the sequence learning problem based on the following instructions:

0. Import necessary modules and set random seed to 42 using `tf.random.set_seed(42)`.
1. Download the data and convert to dataframe as shown in the tutorial. Use the following code to download: (2)

```

!wget https://s3.amazonaws.com/keras-datasets/jena_climate_2009_2016.csv.zip
!unzip jena_climate_2009_2016.csv.zip

```

2. Identify and report the column index for the feature 'Density'. Plot the Density to visualize patterns
 - across multiple days
 - across multiple years.(3)
3. Find the number of datapoints for the train-val-test split of the sequence data based on the following percentages (2)
 - Training set: 55%
 - Validation set: 25%

- Test set : 20%
4. Normalize the raw_data (to mean=0, std=1). Report the mean and std of the training and test sets of the feature '**Density**', pre and post normalization, i.e. compute the following. (5)
- Training set:
 - pre normalization mean =
 - pre normalization std =
 - Test set:
 - post normalization mean =
 - post normalization std = Comment on observed results
5. Use `keras.utils.timeseries_dataset_from_array` to make appropriate datasets from the numbers computed in Q3. (7)
- Use sampling rate = 6,
 - `sequence_length` must contain **6** days of temporal data (Remember: In the raw_data provided, data was recorded every 10 minutes),
 - Modify delay to predict the density **2** days from current time
 - use `batch_size` = 256
 - `Shuffle` = True
 - Info: Use appropriate start and end indices in arguments.
6. Modify the `evaluate_naive_method(dataset)` function to return MAE of the test and validation datasets. (Remember, this our common sense baseline, which assumes that the density changes are purely periodic with a period of 24 hours.) (2)
7. Build the [following model \(click on this link\)](#). Comment on the activation argument in LSTM layer. (3)
8. Compile it using (1)
- `optimizer`= rmsprop with learning rate 0.1
 - `loss`="mse",
 - `metrics`=["mae"]