

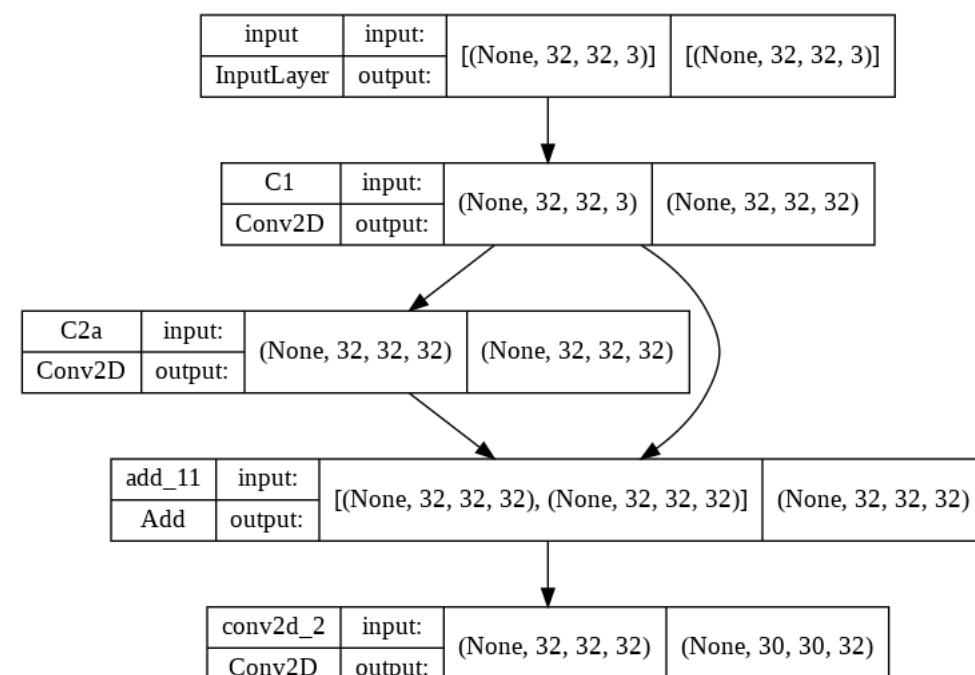
```
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.utils import plot_model
```

Implementing Convolutional Neural Networks

by Sanjeev Gupta

Build model with residual connection

```
inputs = keras.Input(shape=(32, 32, 3), name="input")
x = layers.Conv2D(32, 3, activation="relu", padding="same", name="C1")(inputs) # Q: No. of filters and kernel size? A:
residual = x
x = layers.Conv2D(32, 3, activation="relu", padding="same", name="C2a")(x)
x = layers.add([x, residual]) # Q: Do x and residual have the same shape? A:
x = layers.Conv2D(32, 3)(x)
model = keras.Model(inputs=inputs, outputs=x)
plot_model(model, show_shapes=True)
```



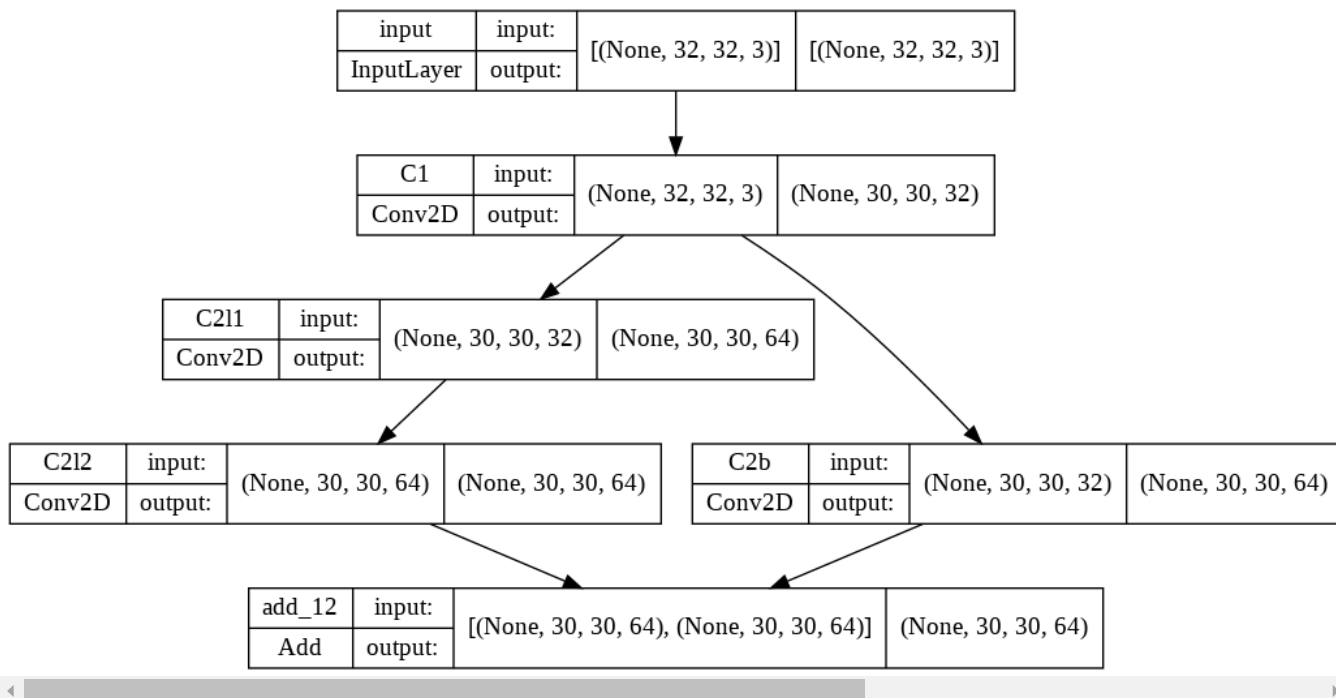
Residual branch may contain 1 layer to make sure addition is possible, i.e. accomodate sizes.

Build model with residual connection - layer in the residual branch

```
inputs = keras.Input(shape=(32, 32, 3), name="input")
x = layers.Conv2D(32, 3, activation="relu", name="C1")(inputs) # Q: No. of filters and kernel size?
residual = x
x = layers.Conv2D(64, 3, activation="relu", padding="same", name="C2l1")(x)
x = layers.Conv2D(64, 3, activation="relu", padding="same", name="C2l2")(x)

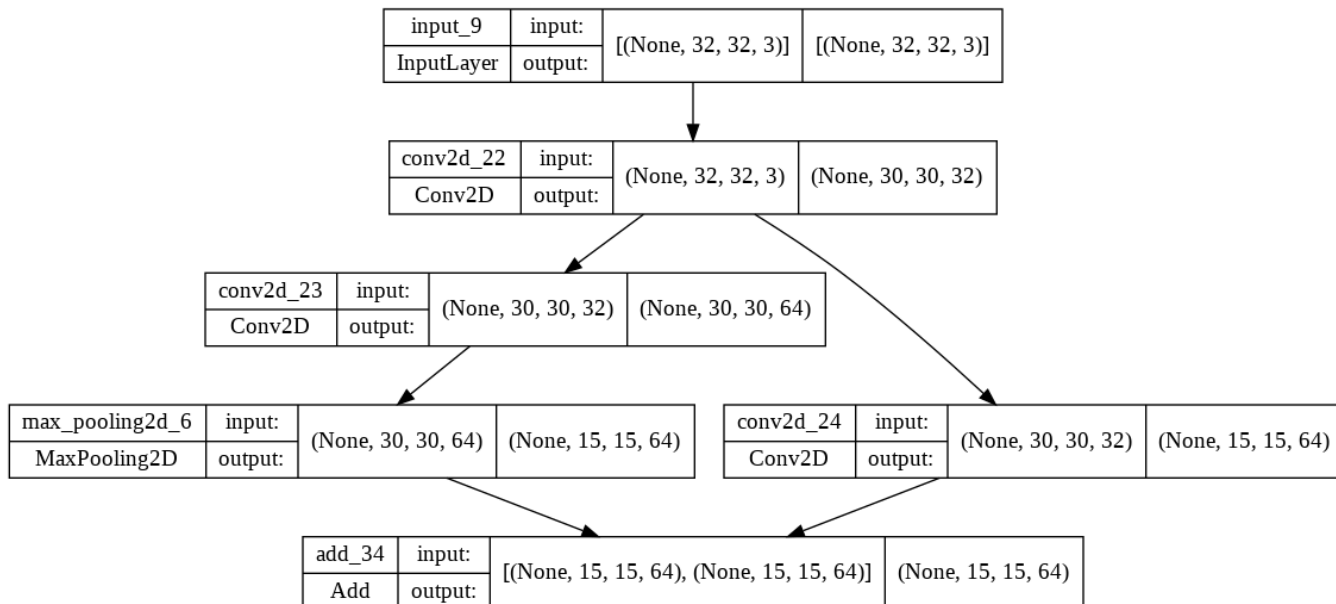
residual = layers.Conv2D(64, 1, name="C2b")(residual) # Q: Why 1? A:
x = layers.add([x, residual]) # Q: Do x and residual have the same shape?
model = keras.Model(inputs=inputs, outputs=x)

plot_model(model, show_shapes=True)
```



Important: Add layers of **same** shape !

```
# model with residual connections and a max pool layer in between
inputs = keras.Input(shape=(32, 32, 3))
x = layers.Conv2D(32, 3, activation="relu")(inputs)
residual = x
x = layers.Conv2D(64, 3, activation="relu", padding="same")(x)
x = layers.MaxPooling2D(2, padding="same")(x) # Max pooling layer reduces the dmsion. #Q: By how much ? A:
residual = layers.Conv2D(64, 1, strides=2)(residual) #Need a stride=2 to accomodate for the maxpool downsampling in the other branch
x = layers.add([x, residual])
model = keras.Model(inputs=inputs, outputs=x)
plot_model(model, show_shapes=True)
```



With residual connections, you can build networks of arbitrary depth, without having to worry about vanishing gradients.

We will see an example later.

Intuitions on why residual blocks work:

- Shorter path for gradients

✓ Batch Normalization

- Adaptively normalize data even as the mean and variance change over time during training
- During training, it uses the mean and variance of the current batch of data to normalize samples
- During inference (when a big enough batch of representative data may not be available), it uses an exponential moving average of the batch-wise mean and variance of the data seen during training.

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots x_m\}$;

Parameters to be learned: γ, β


Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$



input layer

hidden layer 1

hidden layer 2

output layer

Now let's try to calculate the no. of params introduced because of batch normalization

Model: "sequential_3"

Layer (type)	Output Shape	Param #
flatten_3 (Flatten)	(None, 784)	0
batch_normalization_v2 (Batch Normalization)	(None, 784)	3136
dense_50 (Dense)	(None, 300)	235500
batch_normalization_v2_1 (Batch Normalization)	(None, 300)	1200
dense_51 (Dense)	(None, 100)	30100
batch_normalization_v2_2 (Batch Normalization)	(None, 100)	400
dense_52 (Dense)	(None, 10)	1010

Total params: 271,346

Trainable params: 268,978

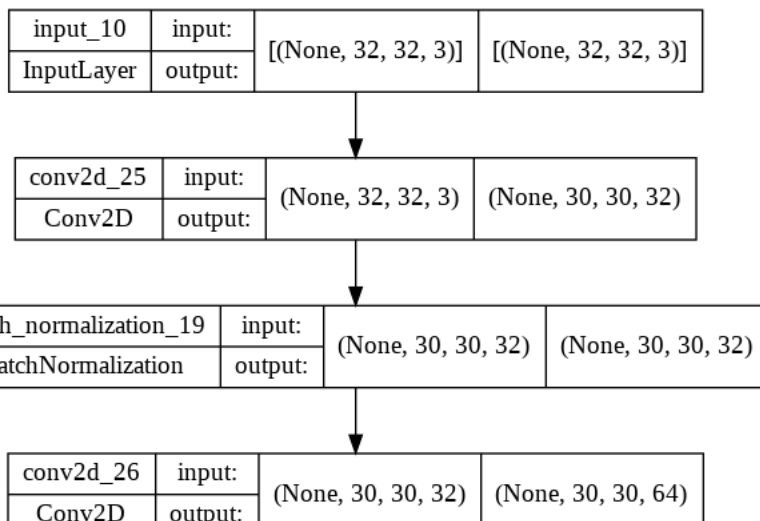
Non-trainable params: 2,368

```

# Model with a batch normalization layer
inputs = keras.Input(shape=(32, 32, 3))
# Because the output of the Conv2D layer gets normalized, the layer doesn't need its own bias vector

```

```
x = layers.Conv2D(32, 3, activation="relu", use_bias=False)(inputs)
x = layers.BatchNormalization()(x)
x = layers.Conv2D(64, 3, activation="relu", padding="same")(x)
model = keras.Model(inputs=inputs, outputs=x)
plot_model(model, show_shapes=True)
```



```
model.summary()
```



Model: "model_17"

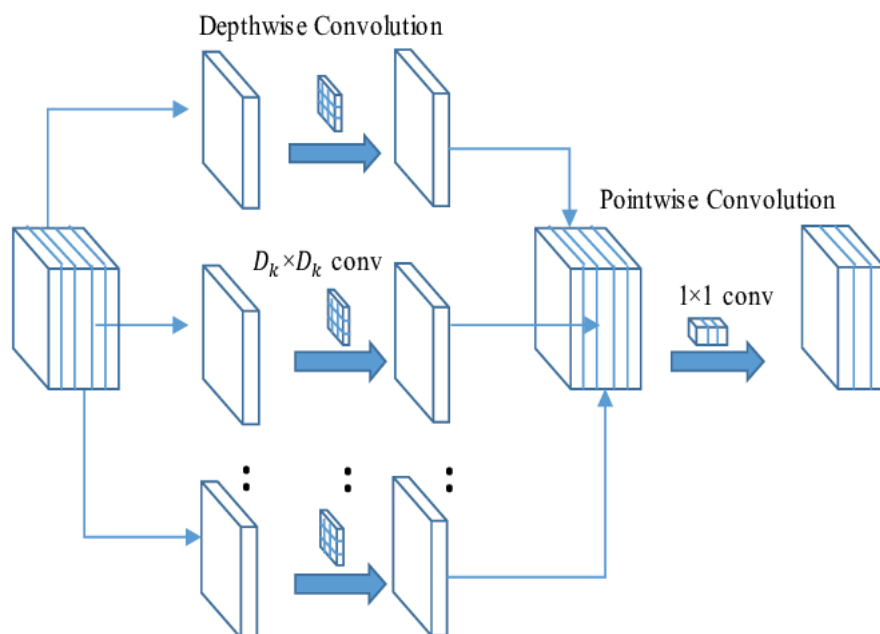

Layer (type)	Output Shape	Param #
input_10 (InputLayer)	[(None, 32, 32, 3)]	0
conv2d_25 (Conv2D)	(None, 30, 30, 32)	864
batch_normalization_19 (BatchNormalization)	(None, 30, 30, 32)	128
conv2d_26 (Conv2D)	(None, 30, 30, 64)	18496
Total params: 19,488		
Trainable params: 19,424		
Non-trainable params: 64		

Intuitions:

- Batch Normalization is also a (weak) regularization method.
 - increases no. of params
 - but also adds noise ~ data augmentation ~ dropout

✓ Depthwise separable convolutions

- This layer performs a spatial convolution on each channel of its input, independently, before mixing output channels via a pointwise convolution
- This makes your model smaller and acts as a strong prior
- We impose a strong prior by assuming that spatial patterns and cross-channel patterns can be modelled separately.
- This is equivalent to separating the learning of spatial features and the learning of channel-wise features.
- depthwise separable convolution relies on the assumption that spatial locations in intermediate activations are highly correlated, but different channels are highly independent.
- So we never use depthwise separable convolution after the input layer. Because RGB channels are **highly corelated**



Let's quickly look at the code first

```
# Building a model with Separable Conv layer
inputs = keras.Input(shape=(32, 32, 3))
x = layers.Conv2D(32, 3, activation="relu")(inputs)
x = layers.SeparableConv2D(64, 3, activation="relu", padding="same")(x)
sep_model = keras.Model(inputs=inputs, outputs=x)
plot_model(sep_model, show_shapes=True)
sep_model.summary()
# Q: Verify the no. of params in the separable_conv2D layer A: 2400 = (32*9) + (32*1*1*64) + 64
```

Model: "model_5"

Layer (type)	Output Shape	Param #
input_3 (InputLayer)	[(None, 32, 32, 3)]	0
conv2d_5 (Conv2D)	(None, 30, 30, 32)	896
separable_conv2d (Separable Conv2D)	(None, 30, 30, 64)	2400
Total params: 3,296		
Trainable params: 3,296		
Non-trainable params: 0		

Let's compare the above with a model where we replace the SeparableConv2D with a Conv2D layer

```
inputs = keras.Input(shape=(32, 32, 3))
x = layers.Conv2D(32, 3, activation="relu")(inputs)
x = layers.Conv2D(64, 3, activation="relu", padding="same")(x)
model = keras.Model(inputs=inputs, outputs=x)
```


```
plot_model(model, show_shapes=True)
```

```
model.summary()
```

```
# Q: Why does sep_model have much less params? A: Depthwise and pointwise convs are done independently
```

```
Model: "model_6"
```

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 32, 32, 3)]	0
conv2d_6 (Conv2D)	(None, 30, 30, 32)	896
conv2d_7 (Conv2D)	(None, 30, 30, 64)	18496
Total params: 19,392		
Trainable params: 19,392		
Non-trainable params: 0		

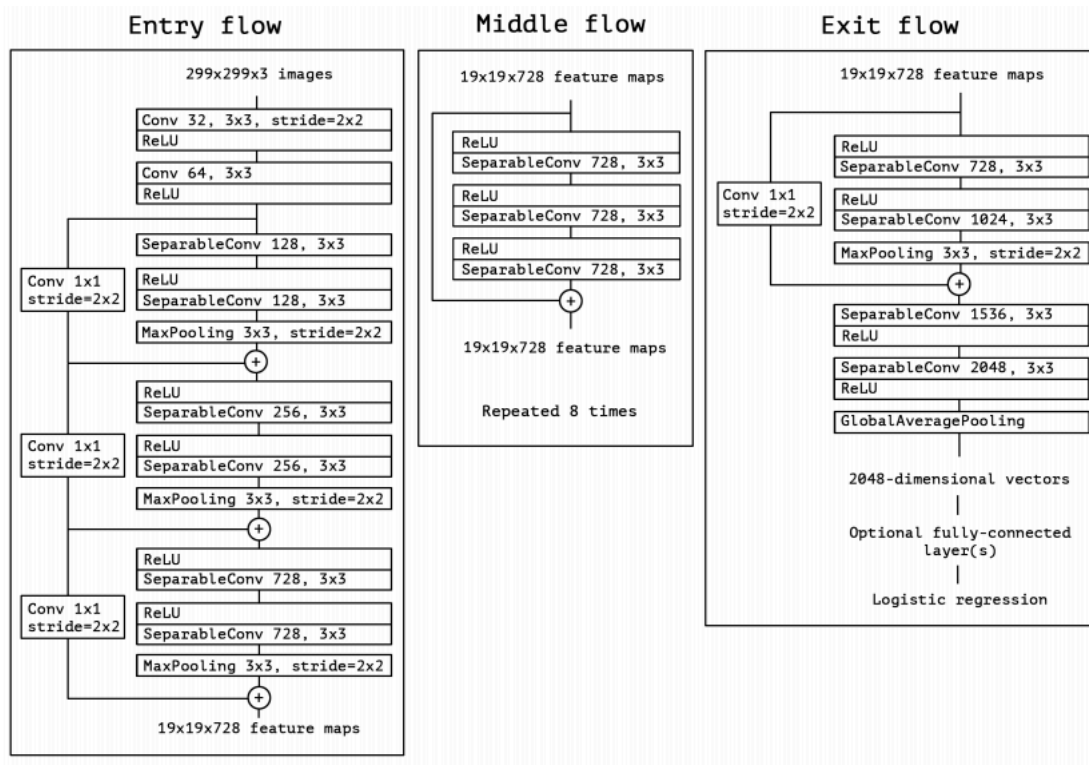


Start coding or [generate](#) with AI.

✓ A mini Xception-like model

We'll build a model like the Xception model, but a smaller version.

But first let's see what the actual Xception model looks like.



Q: In middle- flow blocks, what arguments do you give to the SepConv layer ?

A: HW question

Let's download the cats-vs-dogs data again and create datasets.

Download the dataset

```
!git clone -b master https://github.com/sumanthk94/DA_225-o.git
```

```

Cloning into 'DA_225-o'...
remote: Enumerating objects: 19060, done.
remote: Counting objects: 100% (5024/5024), done.
remote: Compressing objects: 100% (5021/5021), done.
remote: Total 19060 (delta 0), reused 5024 (delta 0), pack-reused 14036
Receiving objects: 100% (19060/19060), 570.44 MiB | 25.44 MiB/s, done.
```

defining path names for futur use

```
data_dir = 'DA_225-o/cats_vs_dogs_small'
```

```
train_path = data_dir + '/train'
```

```
validation_path = data_dir + '/validation'
```

```
test_path = data_dir + '/test'
```

creating datasets using utility

```
from tensorflow.keras.utils import image_dataset_from_directory
```

```

train_dataset = image_dataset_from_directory(
    train_path,
    image_size=(180, 180), # Resize the images to (180,180)
    batch_size=32)

```

```

validation_dataset = image_dataset_from_directory(
    validation_path,
    image_size=(180, 180),
    batch_size=32)

```

```

test_dataset = image_dataset_from_directory(
    test_path,
    image_size=(180, 180),
    batch_size=32)

```

```

Found 2000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.
Found 2000 files belonging to 2 classes.
```

```
import keras
```

```
from keras import layers
```

```
inputs = keras.Input(shape=(180, 180, 3))
```

```
# x = data_augmentation(inputs)
```



```
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=5, use_bias=False)(x) # Q: why not use depth-wise sep conv here? A: RGB channels in input image

for size in [32, 64, 128, 256, 512]:
    residual = x

    x = layers.BatchNormalization()(x) # We can also apply BN just before the activation
    x = layers.Activation("relu")(x)
    x = layers.SeparableConv2D(size, 3, padding="same", use_bias=False)(x)

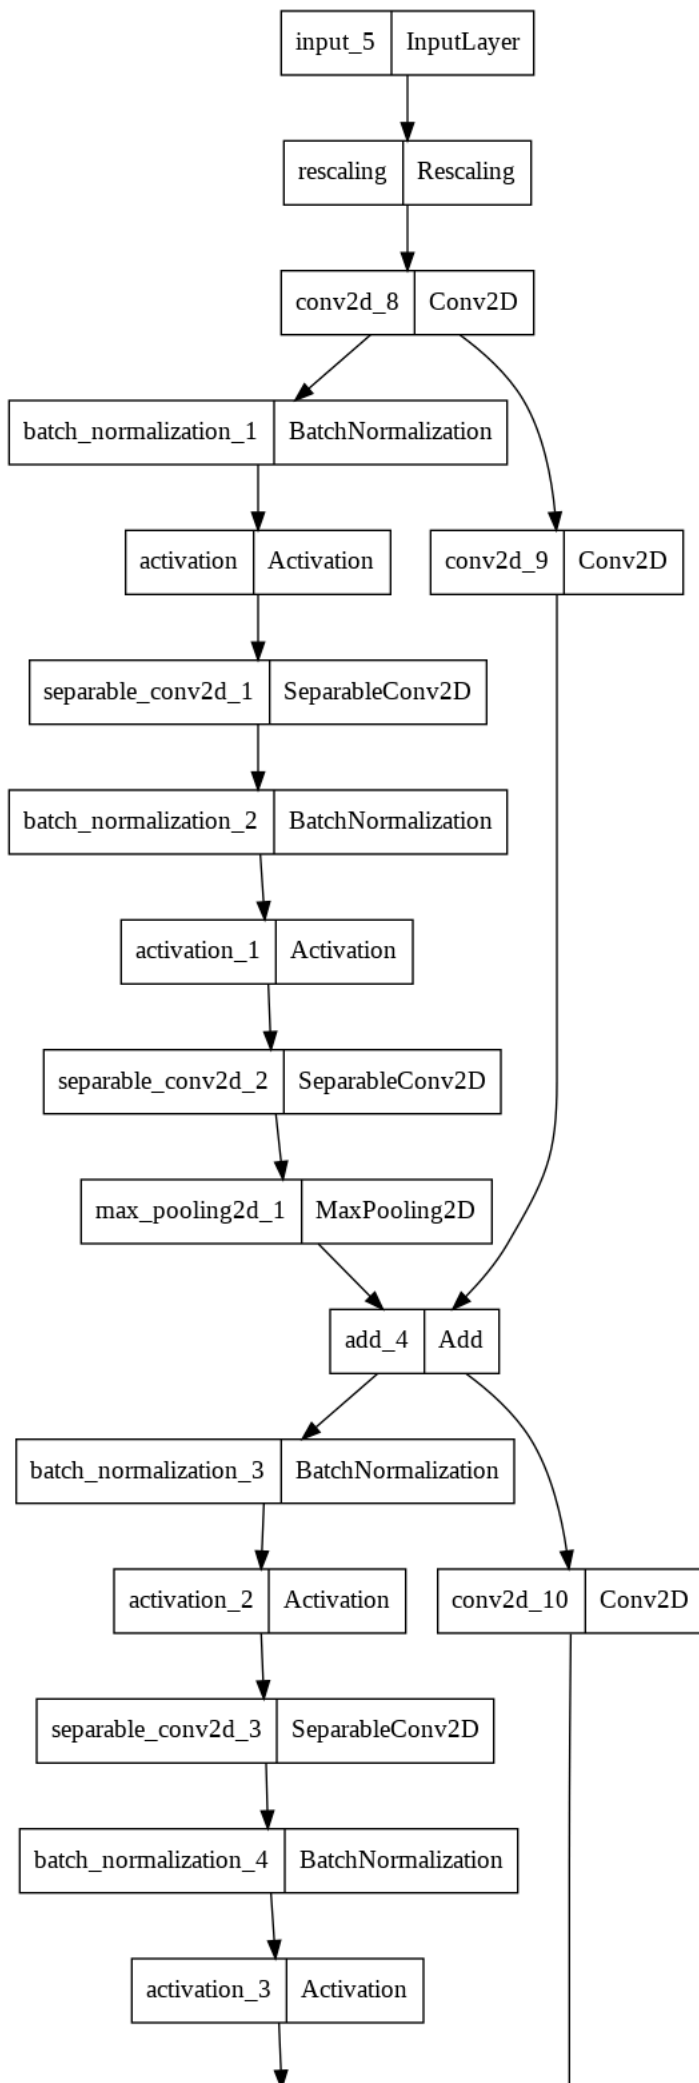
    x = layers.BatchNormalization()(x)
    x = layers.Activation("relu")(x)
    x = layers.SeparableConv2D(size, 3, padding="same", use_bias=False)(x)

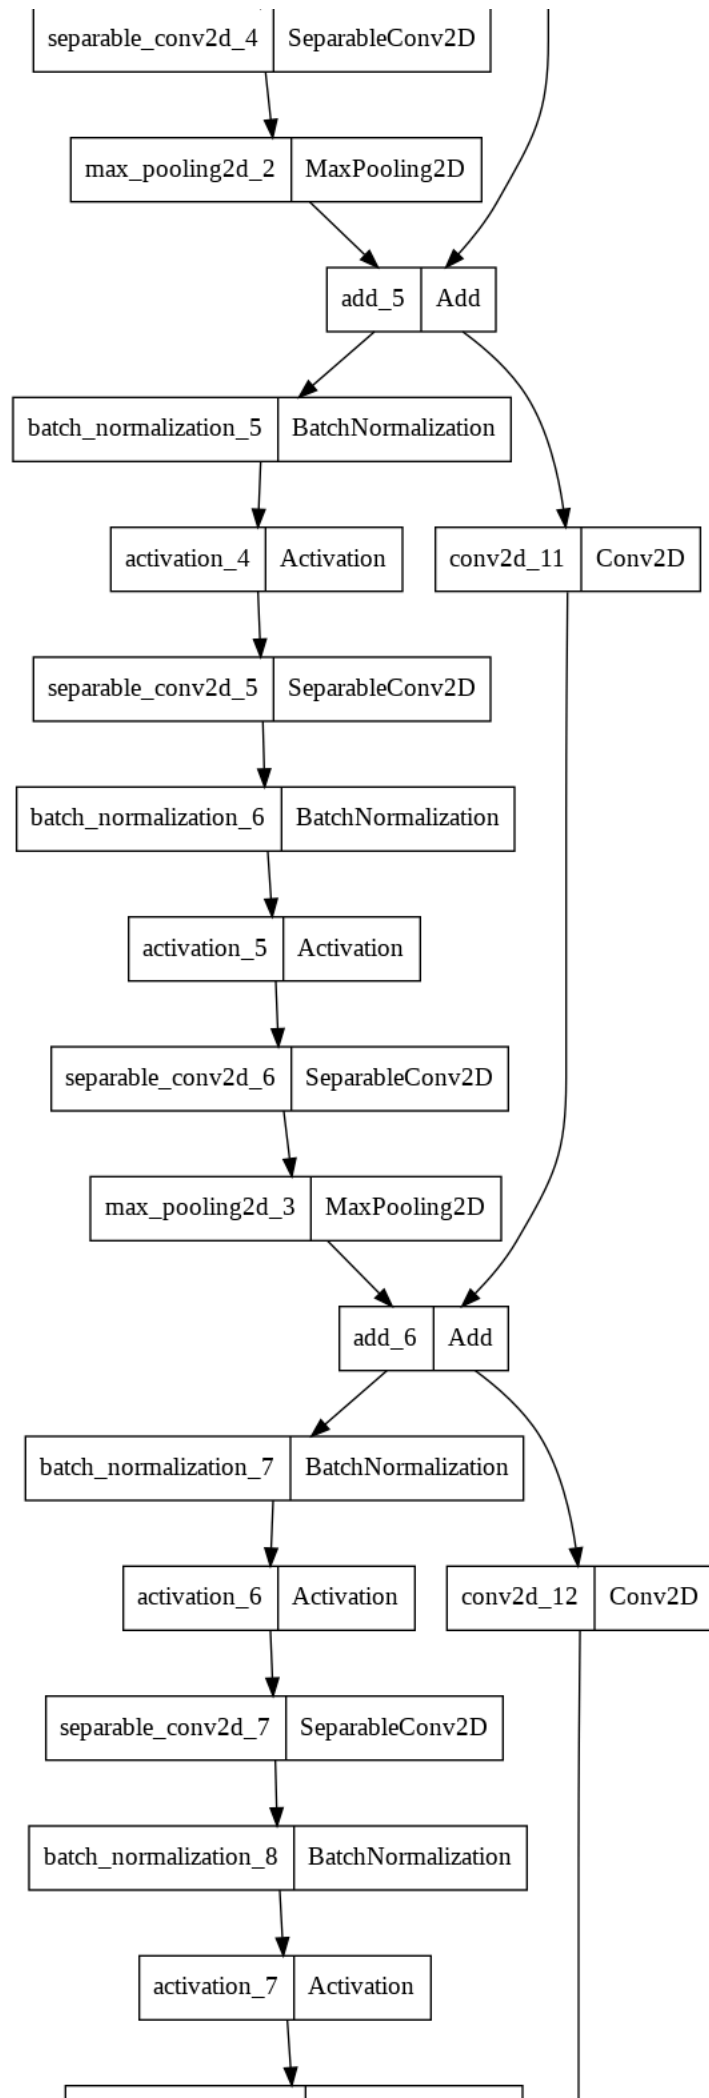
    x = layers.MaxPooling2D(3, strides=2, padding="same")(x)

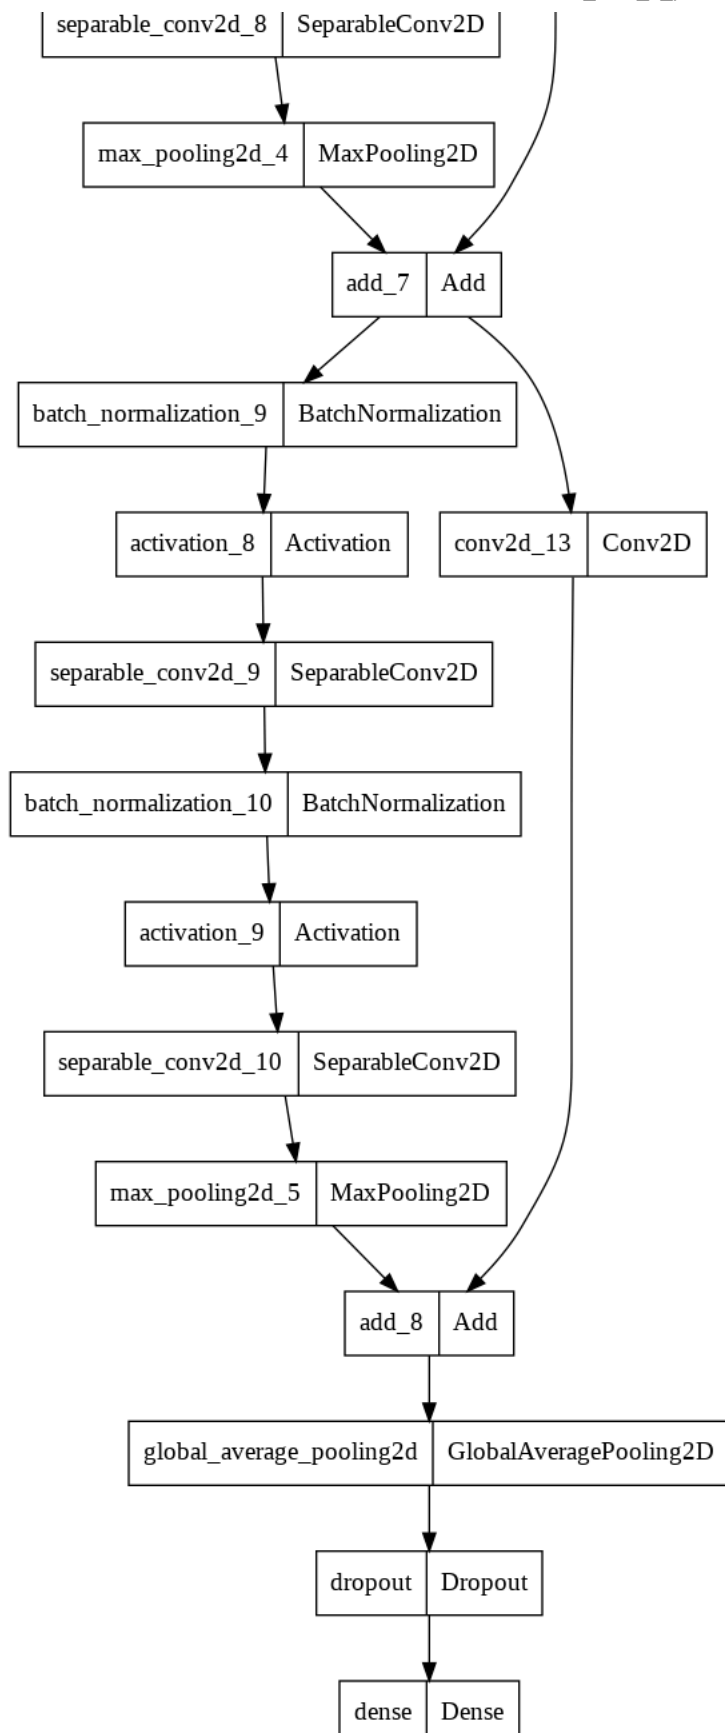
    residual = layers.Conv2D(
        size, 1, strides=2, padding="same", use_bias=False)(residual)
    x = layers.add([x, residual])

x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

plot_model(model)
```



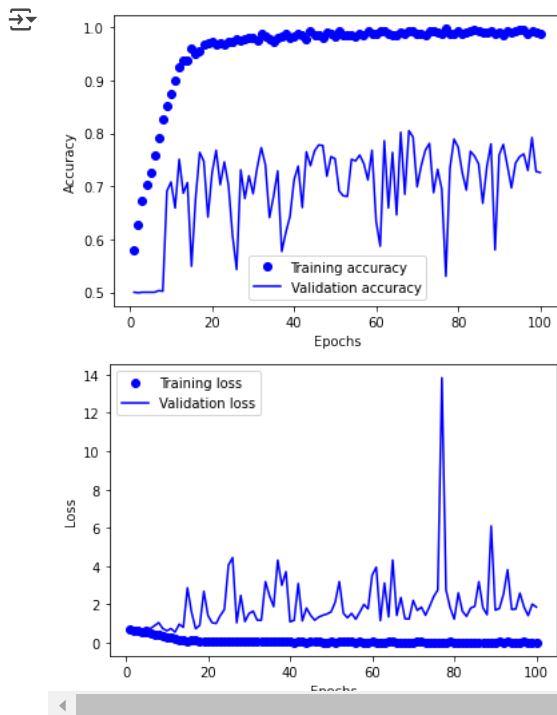




```
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
history = model.fit(train_dataset,
                   epochs=100,
                   validation_data=validation_dataset)
```

```
Epoch 72/100
63/63 [=====] - 16s 254ms/step - loss: 0.0436 - accuracy: 0.9855 - val_loss: 1.8592 - val_accuracy: 0.76
Epoch 73/100
63/63 [=====] - 16s 259ms/step - loss: 0.0213 - accuracy: 0.9930 - val_loss: 1.4377 - val_accuracy: 0.78
Epoch 74/100
63/63 [=====] - 16s 256ms/step - loss: 0.0262 - accuracy: 0.9925 - val_loss: 1.8999 - val_accuracy: 0.68
Epoch 75/100
63/63 [=====] - 16s 259ms/step - loss: 0.0238 - accuracy: 0.9915 - val_loss: 2.4129 - val_accuracy: 0.73
Epoch 76/100
63/63 [=====] - 16s 255ms/step - loss: 0.0348 - accuracy: 0.9890 - val_loss: 2.7587 - val_accuracy: 0.69
Epoch 77/100
63/63 [=====] - 16s 256ms/step - loss: 0.0099 - accuracy: 0.9975 - val_loss: 13.8351 - val_accuracy: 0.5
Epoch 78/100
63/63 [=====] - 16s 261ms/step - loss: 0.0286 - accuracy: 0.9890 - val_loss: 2.7523 - val_accuracy: 0.73
Epoch 79/100
63/63 [=====] - 16s 256ms/step - loss: 0.0458 - accuracy: 0.9885 - val_loss: 1.8347 - val_accuracy: 0.78
Epoch 80/100
63/63 [=====] - 16s 257ms/step - loss: 0.0238 - accuracy: 0.9930 - val_loss: 1.2373 - val_accuracy: 0.77
Epoch 81/100
63/63 [=====] - 16s 255ms/step - loss: 0.0357 - accuracy: 0.9885 - val_loss: 2.6036 - val_accuracy: 0.72
Epoch 82/100
63/63 [=====] - 16s 260ms/step - loss: 0.0276 - accuracy: 0.9915 - val_loss: 1.6421 - val_accuracy: 0.69
Epoch 83/100
63/63 [=====] - 16s 260ms/step - loss: 0.0207 - accuracy: 0.9940 - val_loss: 1.3811 - val_accuracy: 0.76
Epoch 84/100
63/63 [=====] - 16s 260ms/step - loss: 0.0252 - accuracy: 0.9945 - val_loss: 1.8127 - val_accuracy: 0.75
Epoch 85/100
63/63 [=====] - 16s 258ms/step - loss: 0.0171 - accuracy: 0.9930 - val_loss: 1.8998 - val_accuracy: 0.74
Epoch 86/100
63/63 [=====] - 16s 258ms/step - loss: 0.0357 - accuracy: 0.9900 - val_loss: 3.1893 - val_accuracy: 0.66
Epoch 87/100
63/63 [=====] - 16s 258ms/step - loss: 0.0345 - accuracy: 0.9910 - val_loss: 1.8249 - val_accuracy: 0.73
Epoch 88/100
63/63 [=====] - 16s 256ms/step - loss: 0.0132 - accuracy: 0.9935 - val_loss: 1.4656 - val_accuracy: 0.78
Epoch 89/100
63/63 [=====] - 16s 260ms/step - loss: 0.0320 - accuracy: 0.9875 - val_loss: 6.0967 - val_accuracy: 0.58
Epoch 90/100
63/63 [=====] - 16s 258ms/step - loss: 0.0264 - accuracy: 0.9915 - val_loss: 1.7044 - val_accuracy: 0.75
Epoch 91/100
63/63 [=====] - 16s 260ms/step - loss: 0.0490 - accuracy: 0.9865 - val_loss: 1.7897 - val_accuracy: 0.77
Epoch 92/100
63/63 [=====] - 16s 257ms/step - loss: 0.0293 - accuracy: 0.9920 - val_loss: 2.5363 - val_accuracy: 0.73
Epoch 93/100
63/63 [=====] - 16s 259ms/step - loss: 0.0407 - accuracy: 0.9900 - val_loss: 3.8167 - val_accuracy: 0.69
Epoch 94/100
63/63 [=====] - 16s 254ms/step - loss: 0.0279 - accuracy: 0.9920 - val_loss: 1.7427 - val_accuracy: 0.74
Epoch 95/100
63/63 [=====] - 16s 258ms/step - loss: 0.0120 - accuracy: 0.9965 - val_loss: 1.7606 - val_accuracy: 0.75
Epoch 96/100
63/63 [=====] - 16s 257ms/step - loss: 0.0105 - accuracy: 0.9960 - val_loss: 2.5897 - val_accuracy: 0.76
Epoch 97/100
63/63 [=====] - 16s 259ms/step - loss: 0.0490 - accuracy: 0.9885 - val_loss: 1.8709 - val_accuracy: 0.73
Epoch 98/100
63/63 [=====] - 16s 256ms/step - loss: 0.0205 - accuracy: 0.9935 - val_loss: 1.4192 - val_accuracy: 0.79
Epoch 99/100
63/63 [=====] - 16s 256ms/step - loss: 0.0311 - accuracy: 0.9905 - val_loss: 2.0145 - val_accuracy: 0.72
Epoch 100/100
63/63 [=====] - 16s 257ms/step - loss: 0.0270 - accuracy: 0.9900 - val_loss: 1.8610 - val_accuracy: 0.76
```

```
data = pd.DataFrame(history.history)
plt.plot(range(1,len(data)+1),data['accuracy'],'bo',label="Training accuracy")
plt.plot(range(1,len(data)+1),data['val_accuracy'],'b',label="Validation accuracy")
plt.legend()
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.show()
plt.figure()
plt.plot(range(1,len(data)+1),data['loss'],'bo',label="Training loss")
plt.plot(range(1,len(data)+1),data['val_loss'],'b',label="Validation loss")
plt.legend()
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.show()
```




Start coding or [generate](#) with AI.

Visualizing intermediate activations

The output of a layer is called its 'activation'. (It's the output of the activation function)

These activations can be visualized by plotting the feature maps.



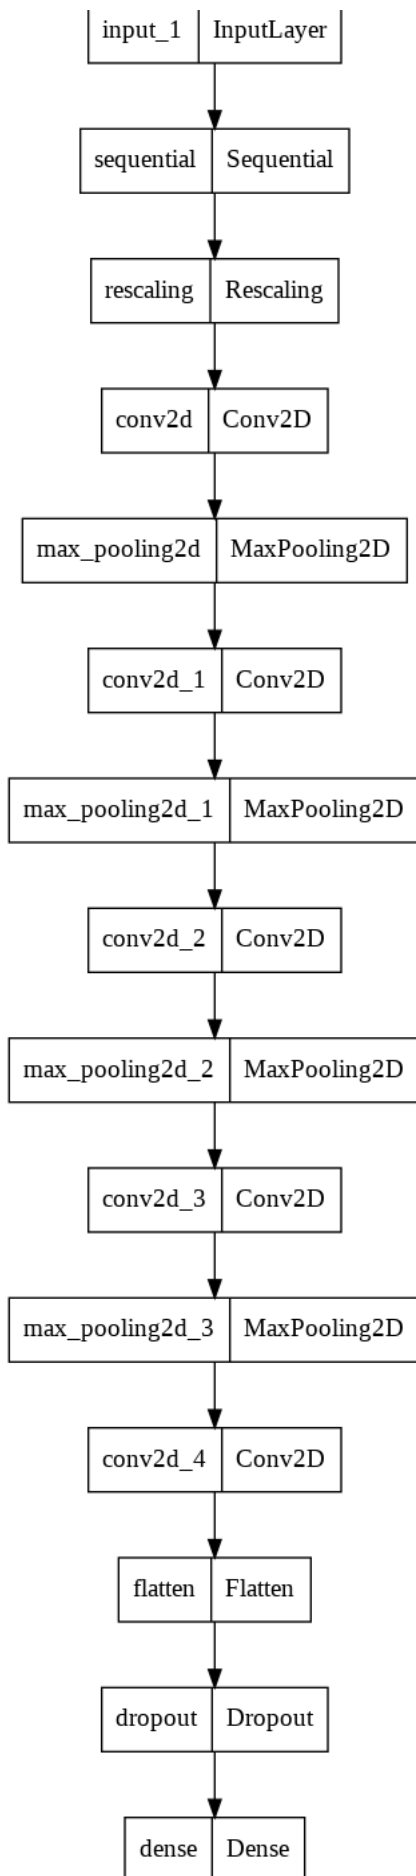
We will plot each feature map independently as a 2D image, since they encode relatively indepent features

```
from tensorflow import keras
from tensorflow.keras.utils import plot_model
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras import layers # <----- Note this
```

```
from google.colab import drive
drive.mount('/content/drive')
%cd "/content/drive/My Drive/"
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
model = keras.models.load_model('/content/drive/MyDrive/Colab Notebooks/convnet_from_scratch_with_augmentation.keras')
plot_model(model)
```



```
# test_loss, test_acc = model.evaluate(test_dataset) #Q: Why don't we have 2 arguments
# print(f"Test accuracy is:{test_acc:.3f}")
```

```
# Preprocessing a single image
```

```
img_path = keras.utils.get_file(fname="cat.jpg",
                                origin="https://img-datasets.s3.amazonaws.com/cat.jpg")
```

```
def get_img_array(img_path, target_size):
    img = keras.utils.load_img(img_path, target_size=target_size)
    array = keras.utils.img_to_array(img) # converts image to np array
    # Add a dimension to transform the array into a "batch" of a single sample.
    array = np.expand_dims(array, axis=0) # Its shape is now (1, 180, 180, 3)
    return array
```

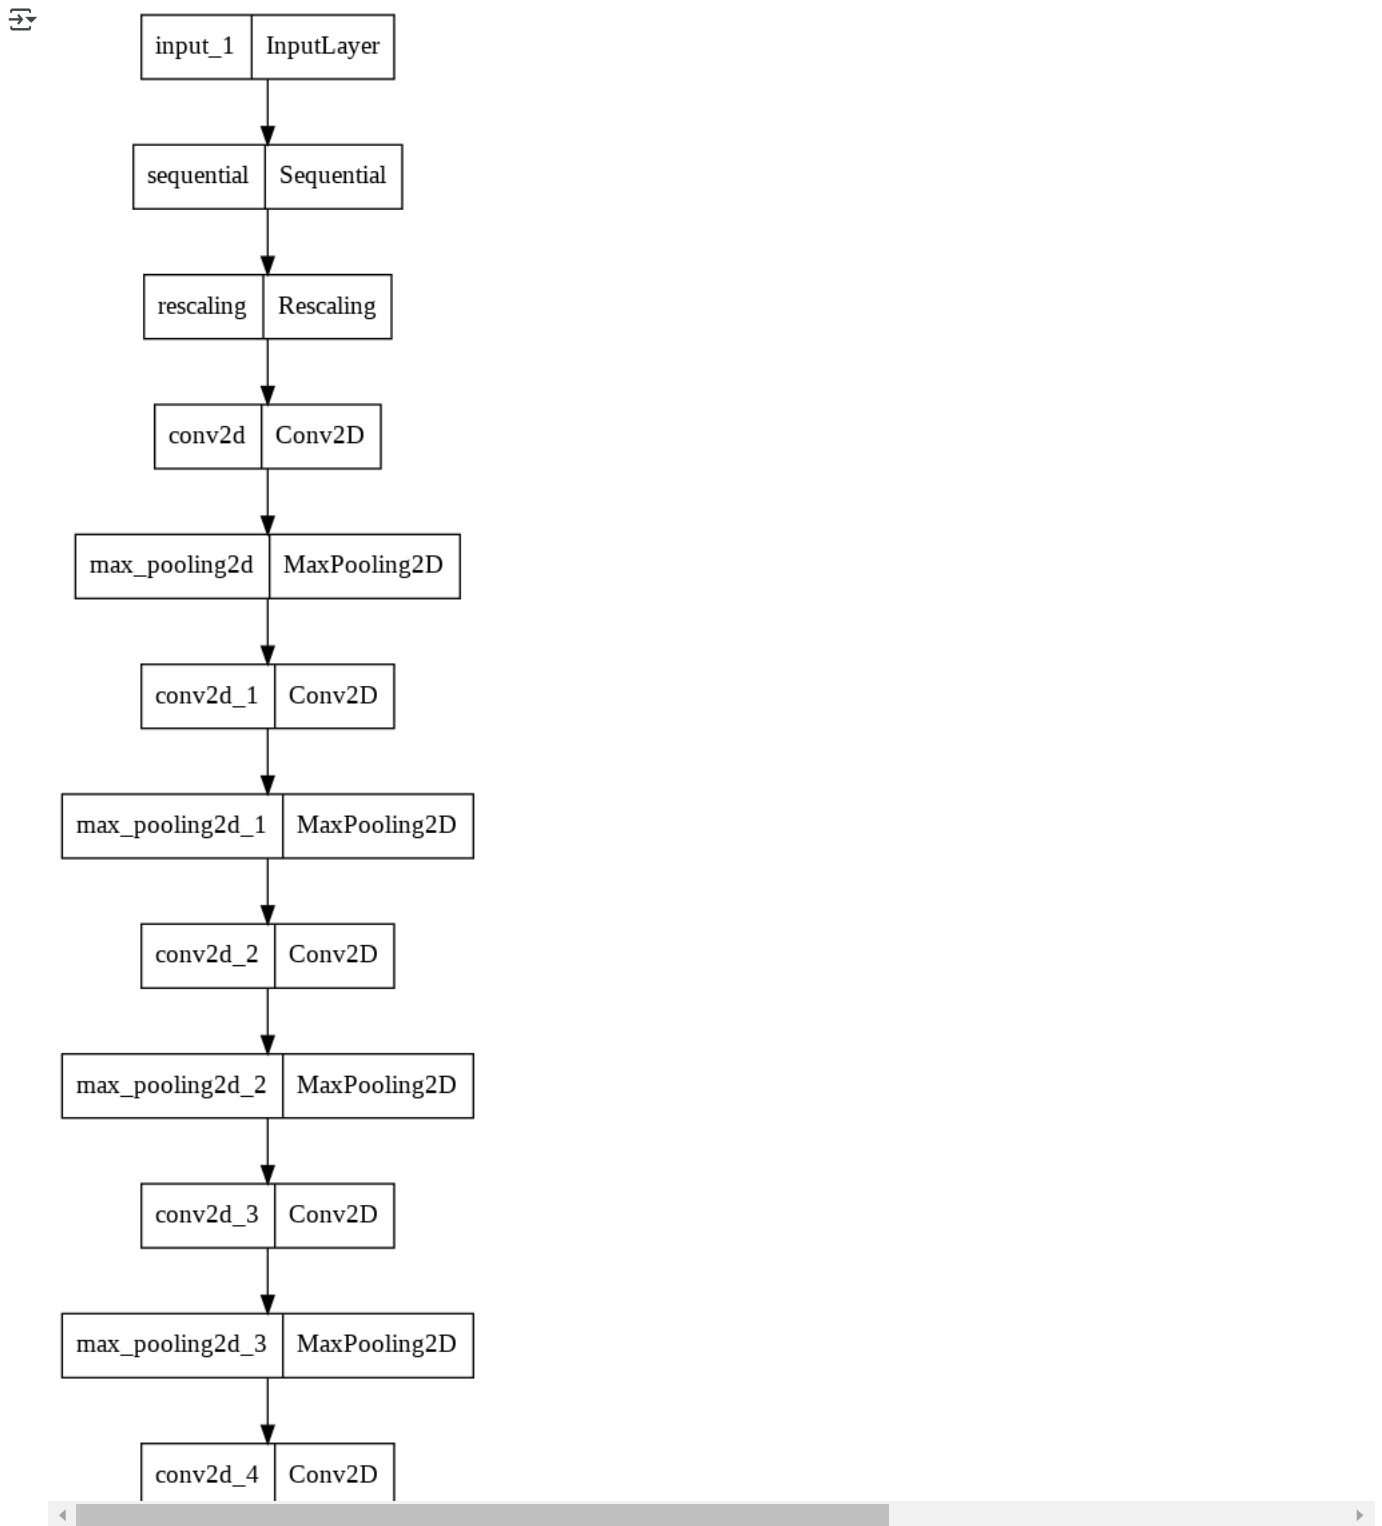
```
img_tensor = get_img_array(img_path, target_size=(180, 180)) #resize image
```

```
plt.axis("off")
plt.imshow(img_tensor[0].astype("uint8"))
plt.show()
```

📄 Downloading data from <https://img-datasets.s3.amazonaws.com/cat.jpg>
 81920/80329 [=====] - 0s 2us/step
 90112/80329 [=====] - 0s 1us/step



```
# Focus here
# Instantiating a model that returns "layer activations"
layer_outputs = []
layer_names = []
for layer in model.layers:
    if isinstance(layer, (layers.Conv2D, layers.MaxPooling2D)): #Q: what does model.layers return ? A: lis of layers of the model
        layer_outputs.append(layer.output)
        # layer_outputs.append(layers.Dense(1)(layer.output)) # To visualise the output in plot_model()
        layer_names.append(layer.name)
activation_model = keras.Model(inputs=model.input, outputs=layer_outputs) #Q: what does layer_outputs contain? A: output of conv2d and r
plot_model(activation_model)
```

```
# Compute layer activations
```


```
activations = activation_model.predict(img_tensor)
print(f"No. of outputs= {len(activations)}")
```

```
first_layer_feature_maps = activations[0] # Q: why is this the first layer's activation? A: index 0
print(f"first_layer_activation.shape= {first_layer_feature_maps.shape}")
```

```
No. of outputs= 9
first_layer_activation.shape= (1, 178, 178, 32)
```

```
# Visualise activation
```


```
import matplotlib.pyplot as plt
plt.matshow(first_layer_feature_maps[0, :, :, 0], cmap="viridis") # Q: which (1st/2nd/..)feature map are we visualizing? A: 1st (0 in the
```




It seems that the filter has detected ____.

Let's look at a feature map after each layer.

```
for i in range(9):  
    plt.matshow(activations[i][0, :, :, 2], cmap="viridis")
```






Note the dimensions on the above images. Successive feature maps are actually of smaller dimensions but scaled to be the same size during visualization.

Now let's visualise all the feature maps of all the layers


```
# Post-processing code - only visualizaton
# Visualizing every channel in every intermediate activation
images_per_row = 16
for layer_name, layer_activation in zip(layer_names, activations):
    n_features = layer_activation.shape[-1]
    size = layer_activation.shape[1]
    n_cols = n_features // images_per_row
    display_grid = np.zeros(((size + 1) * n_cols - 1,
                             images_per_row * (size + 1) - 1))
    for col in range(n_cols):
        for row in range(images_per_row):
            channel_index = col * images_per_row + row
            channel_image = layer_activation[0, :, :, channel_index].copy()
            if channel_image.sum() != 0:
                channel_image -= channel_image.mean()
                channel_image /= channel_image.std()
                channel_image *= 64
                channel_image += 128
            channel_image = np.clip(channel_image, 0, 255).astype("uint8")
            display_grid[
                col * (size + 1): (col + 1) * size + col,
                row * (size + 1) : (row + 1) * size + row] = channel_image
    scale = 1. / size
    plt.figure(figsize=(scale * display_grid.shape[1],
                        scale * display_grid.shape[0]))
    plt.title(layer_name)
    plt.grid(False)
    plt.axis("off")
    plt.imshow(display_grid, aspect="auto", cmap="viridis")
```




conv2d




max_pooling2d




conv2d_1




max_pooling2d_1




conv2d_2




max_pooling2d_2







conv2d_3




max_pooling2d_3





conv2d_4



- The first layer acts as a collection of various edge detectors.
- As you go deeper, the activations become increasingly abstract and less visually interpretable. They begin to encode higher-level concepts such as “cat ear” and “cat eye.”
- The sparsity of the activations increases with the depth of the layer: in the first layer, almost all filters are activated by the input image, but in the following layers, more and more filters are blank. This means the pattern encoded by the filter isn't found in the input image

✓ Visualising convnet filters

- Pick a filter
- Ask the question: What kind of an input image will excite the filter?
- What should the input image be so that you see a (yellow) feature map ?
- In other words, we want to visualize those patterns in the input image that filter picks up and results in high (yellow) values in the feature map.

Instantiating the Xception convolutional base

```
model = keras.applications.xception.Xception(
    weights="imagenet",
    include_top=False)
```

↗ Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/xception/xception_weights_tf_dim_ordering_tf_kern
 83689472/83683744 [=====] - 0s 0us/step
 83697664/83683744 [=====] - 0s 0us/step

Q: Printing the names of conv and sepConv layers in Xception

```
for layer in model.layers:
    if isinstance(layer, (keras.layers.Conv2D, keras.layers.SeparableConv2D)):
        print(layer.name)
```

↗

```
block1_conv1
block1_conv2
block2_sepconv1
block2_sepconv2
conv2d_14
block3_sepconv1
block3_sepconv2
conv2d_15
block4_sepconv1
block4_sepconv2
conv2d_16
block5_sepconv1
block5_sepconv2
block5_sepconv3
block6_sepconv1
block6_sepconv2
block6_sepconv3
block7_sepconv1
block7_sepconv2
block7_sepconv3
block8_sepconv1
block8_sepconv2
block8_sepconv3
block9_sepconv1
block9_sepconv2
block9_sepconv3
block10_sepconv1
block10_sepconv2
block10_sepconv3
block11_sepconv1
block11_sepconv2
block11_sepconv3
block12_sepconv1
block12_sepconv2
block12_sepconv3
block13_sepconv1
block13_sepconv2
conv2d_17
block14_sepconv1
block14_sepconv2
```

Creating a feature extractor model

```
layer_name = "block3_sepconv1"
layer = model.get_layer(name=layer_name)
```