

## Implementing Neural Networks

by **Sanjeev Gupta**

1. Building a neural network from scratch with tensorflow operations
2. Keras Sequential API
3. Keras Functional API
4. Keras Model subclassing
5. Callbacks

### ✓ Import libraries

```
import tensorflow as tf
from tensorflow import keras
import math
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from keras.layers import Dense, Flatten
from keras import Input
from tensorflow.keras.utils import plot_model

from tensorflow.keras.datasets import mnist, fashion_mnist
```

### ✓ 1. Basic Sequential Model

We want to build a sequential model. This means that the layers of our neural network are stacked sequentially. The approach is as follows:

1. First implement a class to build a dense layer. We call it "NaiveDense"
2. Implement a class ("NaiveSequential") to stack the layers sequentially and build or sequential model.

```
# Implementing our dense layer class
class NaiveDense:
    def __init__(self, input_size, output_size, activation): #input and output sizes for the layer
        self.activation = activation

        w_shape = (input_size, output_size) # matrix of weights
        w_initial_value = tf.random.uniform(w_shape, minval=0, maxval=1e-1)
        self.w = tf.Variable(w_initial_value) # we can only update values of tf.Variables

        b_shape = (output_size,) # vector of biases
        b_initial_value = tf.zeros(b_shape)
        self.b = tf.Variable(b_initial_value)

    def __call__(self, inputs): # executed when the class object is used as a function
        return self.activation(tf.matmul(inputs, self.w)+self.b) # self.b is broadcasted

    @property # enables us to use the method as an attribute
    def weights(self):
        return (self.w, self.b)
```

Great! We implemented a dense layer. Now we will stack them together sequentially in our NaiveSequential class

```
class NaiveSequential:
    def __init__(self, layers): # layers: list of layer objects
        self.layers = layers

    def __call__(self, inputs):
        x = inputs
        for layer in self.layers: #output of the prev layer is the input to the next layer
            x = layer(x)
        return x

    @property
    def weights(self):
        weights = []
        for layer in self.layers: # save weights of each layer to a list
            weights += layer.weights # Q: What does layer.weights return?
        return weights # A: layer.weights calls the function layer.weights() since it decorated with @property. It returns (w, b)
```

Great! Now we know how this sequential stacking of dense layer is implemented !

Next, let's instantiate our NaiveSequential class and make our first NN model

```
# define the model
model = NaiveSequential([
    NaiveDense(input_size=28*28,output_size=512,activation=tf.nn.relu),
    NaiveDense(input_size=512,output_size=10,activation=tf.nn.softmax)
])
# Q: What input argument does NaiveSequential take? A: list of layer objects
# Q: What is the input and output dimension of the overall model? A: input dim = 784 , output dim = 10
# Q: Can the output_size of 1 layer be different from the input_size of the next layer? A: No, they have to be the same.
```

28\*28

↩ 784

Now we have our sequential model! But it is untrained and currently not useful.

We must train the model to make it learn useful representations but first we need data !

So let's solve a classification problem by using above sequential model for [MNIST data set](#).


We must

1. Load the data,
2. Reshape the data according to the input shape of the model and
3. Normalize the data

```
# Load data
(train_images, train_labels),(test_images, test_labels) = mnist.load_data()
print(f"train_images.shape = {train_images.shape}")
print(f"train_labels.shape = {train_labels.shape}")
# Q: How many samples does the training set have? A: 60000
```

```
idx = 7
plt.imshow(train_images[idx])
print(f"label for image id {idx}: ",train_labels[idx])
```

↩ train\_images.shape = (60000, 28, 28)  
train\_labels.shape = (60000,)  
label for image id 7: 3



```
# reshape and normalize data
train_images = train_images.reshape((len(train_images),
                                     28*28)).astype("float32")/255
test_images = test_images.reshape((len(test_images),
                                   28*28)).astype("float32")/255

print(train_labels)
```

↩ [5 0 4 ... 5 6 8]

Next we divide the data into batches. For this operation let's implement a class for Batch Generation.

```
class BatchGenerator:

    def __init__(self, images, labels, batch_size=128):
        assert len(images) == len(labels)
        self.index = 0
```

```

self.images = images
self.labels = labels
self.batch_size = batch_size
self.num_batches = math.ceil(len(images)/batch_size)
print(f"batch size = {batch_size}")
print(f"num of batches = {self.num_batches}")

def next(self):
    images = self.images[self.index:self.index + self.batch_size]
    labels = self.labels[self.index:self.index + self.batch_size]
    self.index += self.batch_size
    return images, labels

```

Finally it is time to train the model!


Keras has made life easy for us. Once we have defined the model all we have to do to train it is

1. model.compile()
2. model.fit()

But we should know what goes on behind the scenes. So let's see the steps involved in training a model\

*Training steps:*

1. Compute the predictions using current weights (Forward Pass).
2. Compute the loss value for these predictions.
3. Compute the gradient with respect to model weights.
4. update the weights.



```

# one_training_step function gives the idea of how loss is computed and layer \
# parameters (weights and biases) are updated
def one_training_step(model, images_batch, labels_batch):
    with tf.GradientTape() as tape:          # GradientTape() is the computational graph
        predictions = model(images_batch)    # forward pass.
        per_sample_losses = keras.losses.sparse_categorical_crossentropy( # define loss
            labels_batch, predictions
        )
        average_loss = tf.reduce_mean(per_sample_losses)
    gradients = tape.gradient(average_loss, model.weights) # Compute gradients
    update_weights(gradients, model.weights) # Update the weights
    return average_loss

learning_rate = 1e-3
def update_weights(gradients, weights):
    for g,w in zip(gradients, weights):
        w.assign_sub(g*learning_rate) # w -= g*lr

```

```
# Full training loop
def fit(model, images, labels, epochs, batch_size=128):
    for epoch in range(epochs):
        # repeat for epochs
        print(f"Epoch {epoch}")
        batch_generator = BatchGenerator(images, labels)
        for batch_counter in range(batch_generator.num_batches):
            # go through all mini-batches in the data
            images_batch, labels_batch = batch_generator.next()
            loss = one_training_step(model, images_batch, labels_batch)
            if batch_counter%100 == 0:
                print(f"loss at batch {batch_counter}:{loss:.2f}")

# Q: Identify the 4 training steps in the above code
# A: They are present in the one_training_step function. See comments.
```

Now let's train the model on MNIST data set.

```
fit(model, train_images, train_labels, epochs=10, batch_size=128)
```

```
# Q: we didn't do a compile step.... or did we? (read again after seeing keras API)
# A: In model.compile(), we pass information about the loss function, optimizer, and evaluation metric. \
#     In our naive implementation, instead of defining a separate compile() function, we have defined the loss \
#     inside one_training_step; implemented the optimizer 'mini-batch gradient descent' in update_weights(); \
#     and we are doing the evaluation (accuracy) separately in a later cell.
```

```
Epoch 0
batch size = 128
num of batches = 469
loss at batch 0:5.57
loss at batch 100:2.24
loss at batch 200:2.21
loss at batch 300:2.11
loss at batch 400:2.22
Epoch 1
batch size = 128
num of batches = 469
loss at batch 0:1.91
loss at batch 100:1.88
loss at batch 200:1.84
loss at batch 300:1.74
loss at batch 400:1.84
Epoch 2
batch size = 128
num of batches = 469
loss at batch 0:1.59
loss at batch 100:1.58
loss at batch 200:1.52
loss at batch 300:1.45
loss at batch 400:1.52
Epoch 3
batch size = 128
num of batches = 469
loss at batch 0:1.33
loss at batch 100:1.34
loss at batch 200:1.25
loss at batch 300:1.23
loss at batch 400:1.29
Epoch 4
batch size = 128
num of batches = 469
loss at batch 0:1.13
loss at batch 100:1.15
loss at batch 200:1.05
loss at batch 300:1.06
loss at batch 400:1.12
Epoch 5
batch size = 128
num of batches = 469
loss at batch 0:0.98
loss at batch 100:1.01
loss at batch 200:0.91
loss at batch 300:0.94
loss at batch 400:1.00
Epoch 6
batch size = 128
num of batches = 469
loss at batch 0:0.87
loss at batch 100:0.91
loss at batch 200:0.80
loss at batch 300:0.85
loss at batch 400:0.91
Epoch 7
batch size = 128
```

After 10 epochs the loss has come down to 0.73!

Our model has definitely learned something. Lets evaluate how accurately it can predict labels for images **it has not seen before**. These are the **images in the test set**.

Lets use the "accuracy" metric. Here we simply find the fraction of times the model succeeded in predicting the correct label.

When using Keras, we would mention this metric in `model.compile()`. (More on Keras later)

```
# Evaluation step
predictions = model(test_images).numpy()
predicted_labels = np.argmax(predictions,axis=1)
matches = predicted_labels == test_labels
print(f"accuracy:{matches.mean():.2f}")
print(predictions.shape)
print(predicted_labels)
print(matches)
```


```
↗ accuracy:0.82
(10000, 10)
[7 2 1 ... 4 8 6]
[ True  True  True ...  True False  True]
```

```
idx = 0
plt.imshow(test_images[idx].reshape(28,28))
print(predicted_labels[idx], test_labels[idx])
```



## Summary

# The BIG Picture



## ▼ Different APIs

1. Sequential Model
2. Functional API

## 3. Model subclassing

## Model building: from **simple** to **arbitrarily flexible**



### 2. Sequential API

Whatever we have implemented so far can be done alternatively using Sequential class in keras. In the following approach layer are passed as a list

```
# from keras.layers import Dense, Flatten
# from keras import Input

seq_model = keras.Sequential([
    Dense(64, activation="relu"),
    Dense(10, activation="softmax")
])
# Q: Do you notice a difference in arguments of the Dense layers, compared to our implementation?
# A: Yes. No need to mention input size for keras.Sequential. It is inferred by keras itself.
```

Alternatively, instead of passing layers as list, we can build a sequential model by adding layers incrementally to the model.

```
seq_model_inc = keras.Sequential()
seq_model_inc.add(Dense(64, activation="relu"))
seq_model_inc.add(Dense(10, activation="softmax"))
```

Notice that we have not yet provided information of input dimensions.

These layers are referred to as symbolic layers.

Unless until you build the model layer weights are not created.

```
try:
    seq_model_inc.weights
except:
    print("seq_model_inc.weights did not work because model was not built and weights were not initialized.")
```

seq\_model\_inc.weights did not work because model was not built and weights were not initialized.


To create a weights you need to call on some data or call its build method with input shape

```
seq_model_inc.build(input_shape=(None, 3)) #None means it can take any batch size; 3 is the number of features in your input
seq_model.build(input_shape=(None, 3))

print(type(seq_model_inc.weights))
seq_model_inc.weights
```

<class 'list'>

```
[<tf.Variable 'dense_15/kernel:0' shape=(3, 64) dtype=float32, numpy=
array([[ 1.44846976e-01, -2.18670100e-01,  2.29430258e-01,
        -1.12742782e-01,  1.36850089e-01,  1.01428717e-01,
        -1.97391719e-01, -2.33028114e-01, -1.45720556e-01,
         2.20789492e-01, -8.29304010e-02, -2.48935312e-01,
         6.46039546e-02, -1.67079866e-02,  1.26166314e-01,
         1.20301038e-01,  2.29161739e-01, -3.97650599e-02,
        -2.61659771e-01,  1.77647531e-01,  1.50566190e-01,
         1.51140183e-01,  1.02343738e-01,  9.82773006e-02,
        -2.41253480e-01,  2.93537140e-01,  2.20617294e-01,
        -1.70859471e-01, -2.71017343e-01,  7.28799105e-02,
        -2.69627869e-01,  1.87116683e-01,  2.83798873e-01,
        -5.05889654e-02,  2.56579220e-01, -8.36384445e-02,
         2.35531926e-01,  2.91702092e-01,  1.91416949e-01,
        -1.45027846e-01, -1.42255217e-01, -2.58929133e-02,
        -2.61746168e-01, -7.54176527e-02, -1.28255069e-01,
        -2.36875176e-01, -2.05620587e-01,  1.02031231e-02,
         8.78518820e-02, -7.40671903e-02,  1.67623580e-01,
        -5.38714528e-02, -1.76747844e-01, -2.11665481e-01,
         8.26619864e-02, -9.00956839e-02,  2.38787055e-01,
        -2.89154440e-01, -2.00049490e-01, -2.38261253e-01,
        -8.48178267e-02, -3.04153264e-02,  2.77565122e-02,
         9.41936672e-02],
 [ 3.28582525e-03, -1.53826490e-01, -2.07747310e-01,
   -2.85859972e-01, -1.96267068e-01,  1.04508013e-01,
   -2.46805385e-01,  2.88513303e-02, -2.62435108e-01,
    1.49608076e-01,  1.64378703e-01, -2.65429258e-01,
   -1.61449030e-01,  2.38224804e-01, -1.97246030e-01,
    1.37537569e-01,  7.95944035e-02, -1.44790545e-01,
   -2.10254103e-01, -4.22446728e-02, -1.66679859e-01,
    1.38105214e-01,  2.66237855e-02, -2.02402458e-01,
    2.35694528e-01,  1.61121696e-01,  1.64229870e-01,
    2.08356559e-01, -2.50580221e-01,  1.11113548e-01,
    1.86853856e-01, -8.17737132e-02,  1.69566095e-01,
   -1.15439638e-01,  2.46289253e-01, -1.89692974e-01,
   -2.54041076e-01,  7.88029432e-02, -1.33986980e-01,
   -1.26948759e-01,  6.54541254e-02,  8.13392699e-02,
    2.53158867e-01, -2.55190998e-01,  2.17753053e-01,
    1.03417307e-01, -2.14490712e-01,  2.62766302e-01,
    2.93517709e-01, -9.50333476e-03, -2.23486334e-01,
    2.00392246e-01,  1.04625732e-01, -1.44753873e-01,
    4.41914499e-02,  3.47997844e-02,  1.69060737e-01,
    2.31785297e-01, -3.53853405e-02, -1.66206911e-01,
   -1.33526579e-01, -2.68126845e-01, -2.46547043e-01,
    1.63272232e-01],
 [ 1.03941858e-02, -1.73210725e-01,  2.63023078e-01,
   1.94167882e-01,  1.41594261e-01, -1.69026345e-01,
   1.38517261e-01, -5.26740998e-02,  2.46134996e-02,
   4.81922626e-02,  2.74699211e-01, -2.86230117e-01,
   -1.43859819e-01,  2.02106476e-01, -8.32575560e-03,
    2.46754348e-01,  8.95386636e-02,  1.03722185e-01,
   -5.42291850e-02, -9.46594179e-02,  2.53973305e-01,
    2.88602710e-01, -1.62680984e-01, -9.43397880e-02,
    8.43972266e-02,  2.73092330e-01, -2.59637833e-04,
    1.68687165e-01,  9.12820399e-02,  1.61256403e-01,
    7.51611590e-02, -1.39940783e-01, -2.11833864e-01,
   -1.83083535e-01,  2.18872451e-01, -2.84347110e-01]
```



```
seq_model.summary()
```

```
↔ Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
dense_13 (Dense)	(None, 64)	256
dense_14 (Dense)	(None, 10)	650
Total params: 906		
Trainable params: 906		
Non-trainable params: 0		

Lets calculate the no. of params in each layer. For layer 1:

```
(64*3) + 64
```

```
↔ 256
```

Q: Can you verify the number of parameters by a quick caluclation?

A:  $650 = 10 \times 64 + 10$

weight matrix has  $64 \times 10$  weights and 10 biases for the 10 neurons

```
256 + 650
```

```
↔ 906
```

### Specifying input shape in advance

```
model_seq = keras.Sequential(name="sequential_model")
model_seq.add(keras.Input(shape=(3,))) #specifying the input here
model_seq.add(keras.layers.Dense(64, activation=tf.nn.relu, name="first_layer"))
model_seq.add(keras.layers.Dense(10, activation=tf.nn.softmax,
                                name="second_layer"))
```

```
model_seq.summary()
```

```
↔ Model: "sequential_model"
```

Layer (type)	Output Shape	Param #
--------------	--------------	---------



```

=====
first_layer (Dense)          (None, 64)          256

second_layer (Dense)         (None, 10)          650

=====
Total params: 906
Trainable params: 906
Non-trainable params: 0
=====

```

Double-click (or enter) to edit

### ✓ 3. Functional API

Next we will use Keras functional API to create the same model. Keras functional API can create more flexible models than Sequential API. It can handle models with non-linear topology, shared layers, and even multiple inputs or outputs.

Key Idea- Expresses each layer as a function of the previous layer.

```

      (input: 3-dimensional vectors)
            ↓
      [Dense (64 units, relu activation)]
            ↓
      (output: 10 units, softmax activation)

```

```

inputs = Input(shape=(3,), name="input_layer")
features = Dense(64, activation="relu", name="first_layer")(inputs) #f(inputs)
outputs = Dense(10, activation="softmax",
               name="output_layer")(features) #f(features)
fun_model = keras.Model(inputs, outputs)

```

```
fun_model.summary()
```



Model: "model"

```

Layer (type)                 Output Shape          Param #
=====
input_layer (InputLayer)     [(None, 3)]           0

first_layer (Dense)          (None, 64)            256

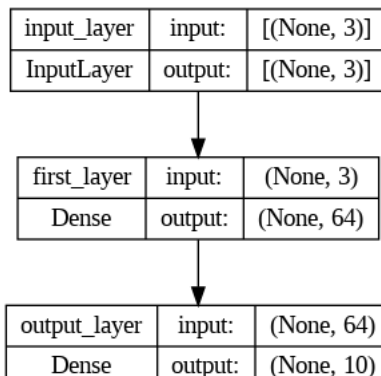
output_layer (Dense)         (None, 10)            650

=====
Total params: 906
Trainable params: 906
Non-trainable params: 0
=====

```

We get effectively the same summary. Because we have implemented the same model using the functional API.

```
plot_model(fun_model, show_shapes=True)
```



Lets see a deeper network.

```

# from keras.layers import Dense
# import keras
# node = Layer(nodes, extra_params)(prev_node)

```

```
inputs = keras.Input(shape=(64,))
dense1 = Dense(32, activation='relu')(inputs)
dense2 = Dense(32, activation='relu')(dense1) #defining dense2 node whose parent is dense1
outputs = Dense(4, activation='softmax')(dense2) #defining output node where parent is dense2
model = keras.Model(inputs=inputs, outputs=outputs, name="linear_topology")
```

```
model.summary()
```

```
Model: "linear_topology"
```

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 64)]	0
dense_4 (Dense)	(None, 32)	2080
dense_5 (Dense)	(None, 32)	1056
dense_6 (Dense)	(None, 4)	132
Total params: 3,268		
Trainable params: 3,268		
Non-trainable params: 0		

```
32*64 + 32
```

```
2080
```

```
plot_model(model)
```



Now let's see an example where the Sequential API would not be sufficient.

**Multi-Input and Multi-output:** Consider an example of building a system to rank customer tickets by priority and route them to the appropriate departments.

Outputs: model need to give two outputs

1. First task of the model is to classify the tickets into priority and non priority (Binary classification)
2. Second task is to route the ticket to appropriate department (Multi-class classification based on the number of departments)

These two task are need to be done simultaneously

Inputs:

1. Title of the ticket (text input)
2. The text body of the ticket (text input)
3. Any tags added by the user

Q. Is it possible to build the model sequentially?

A: No, we cannot build a multi-input , multi-output model through the sequential API, because, by definition itself, the required model is not sequential.

```
vocabulary_size = 10000
num_tags = 100
num_departments = 4
```

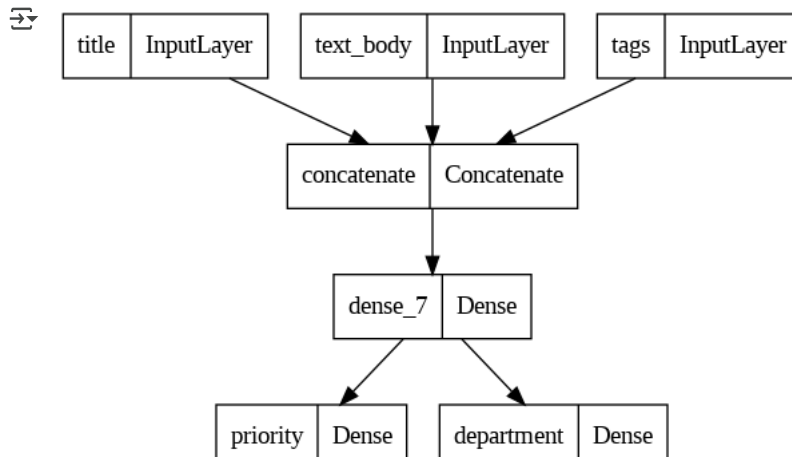
```
# Inputs
title = keras.Input(shape=(vocabulary_size,), name="title")
text_body = keras.Input(shape=(vocabulary_size, ), name="text_body")
tags = keras.Input(shape=(num_tags, ), name="tags")

features = keras.layers.Concatenate()([title, text_body, tags])
features = keras.layers.Dense(64, activation="relu")(features)

# Outputs
priority = keras.layers.Dense(1, activation="sigmoid", # sigmoid for binary classification
                             name="priority")(features)
department = keras.layers.Dense(num_departments, activation="softmax", # softmax for multi-class classification
                                name="department")(features)

model = keras.Model(inputs=[title, text_body, tags],
                    outputs=[priority, department])
```

```
plot_model(model)
```



```
model.summary()
```

```
Model: "model_1"
```

Layer (type)	Output Shape	Param #	Connected to
title (InputLayer)	[(None, 10000)]	0	[]
text_body (InputLayer)	[(None, 10000)]	0	[]
tags (InputLayer)	[(None, 100)]	0	[]
concatenate (Concatenate)	(None, 20100)	0	['title[0][0]', 'text_body[0][0]', 'tags[0][0]']
dense_7 (Dense)	(None, 64)	1286464	['concatenate[0][0]']
priority (Dense)	(None, 1)	65	['dense_7[0][0]']
department (Dense)	(None, 4)	260	['dense_7[0][0]']

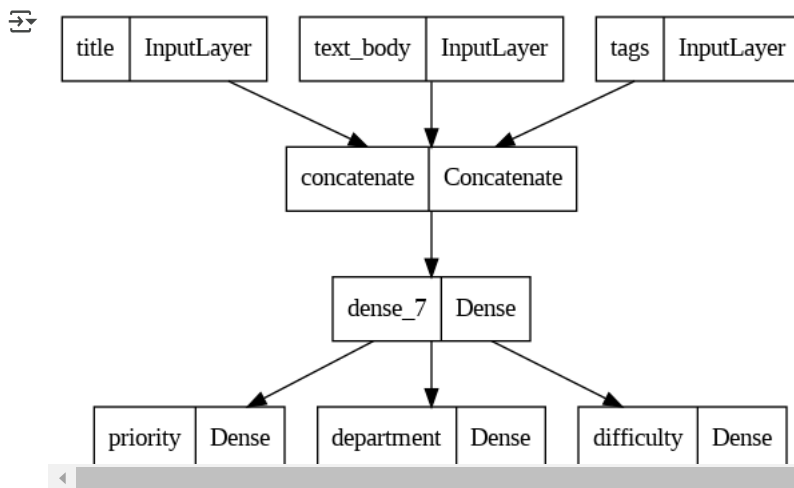
=====  
 Total params: 1,286,789  
 Trainable params: 1,286,789  
 Non-trainable params: 0  
 =====

Reusing the model by training intermediate layer output

```
features = model.layers[4].output
difficulty = keras.layers.Dense(3, activation="softmax",
                                name="difficulty")(features)

new_model = keras.Model(inputs=[title, text_body, tags],
                        outputs=[priority, department, difficulty])

keras.utils.plot_model(new_model)
```



## 4. Subclassing the Model class

We saw how the functional API enabled us to make more complex models compared to the sequential API. We moved up the ladder of progressive disclosure of complexity.

Now let's move a step further.

But first, let us see how to define the same old model by subclassing the Model class.

```

class CustomerTicketModel(keras.Model):

# Define the layers in in the __init__ method
def __init__(self, num_departments):
    super().__init__()
    self.concat_layer = keras.layers.Concatenate()
    self.mixing_layer = keras.layers.Dense(64, activation="relu")
    self.priority_scorer = keras.layers.Dense(1, activation="sigmoid")
    self.department_classifier = keras.layers.Dense(num_departments,
                                                    activation="softmax")

# Define the relationship between layers in in the call method
# See Section 7.2.3 in Francois chollet for more details
def call(self,inputs):
    # input should be dictionary type
    title = inputs["title"]
    text_body = inputs["text_body"]
    tags = inputs["tags"]

    features = self.concat_layer([title, text_body, tags])
    features = self.mixing_layer(features)

    priority = self.priority_scorer(features)
    department = self.department_classifier(features)

    return priority, department

sub_class_model = CustomerTicketModel(num_departments=4)

try:
    sub_class_model.summary()
except:
    print("summary() did not work because we have not built the model")

summary() did not work because we have not built the model

# here model is built by calling the data since build() method is not
# defined in model subclass
# generate random data
title_data = np.random.randint(0, 2, size=(1000,vocabulary_size))
text_body_data = np.random.randint(0, 2, size=(1000,vocabulary_size))
tags_data = np.random.randint(0, 2, size=(1000,num_tags))

priority, department = sub_class_model({"title":title_data,
                                       "text_body":text_body_data,
                                       "tags":tags_data})

```

```
sub_class_model.summary()
```

Model: "customer\_ticket\_model"

Layer (type)	Output Shape	Param #
=====		
concatenate_1 (Concatenate)	multiple	0
dense_8 (Dense)	multiple	1286464
dense_9 (Dense)	multiple	65
dense_10 (Dense)	multiple	260
=====		
Total params: 1,286,789		
Trainable params: 1,286,789		
Non-trainable params: 0		

```
from keras.initializers import RandomNormal
```

```
# Creating a custom LAYER, by subclassing keras.layers.Layer
```

```
class Custom_Dense(keras.layers.Layer):
```

```
    def __init__(self, units, activation=None):
```

```
        super().__init__()
```

```
        self.units = units
```

```
        self.activation = activation
```

```
# Subclassing gives us the flexibility here to initialize weights on our own
```

```
    def build(self, input_shape):
```

```
        input_dim = input_shape[-1]
```

```
        std_dev = np.sqrt(2/(input_dim + self.units))
```

```
        self.W = self.add_weight(shape=(input_dim, self.units),
                                initializer=RandomNormal(stddev=std_dev))
```

```
        self.b = self.add_weight(shape=(self.units,),
                                initializer="zeros")
```

```
    def call(self, inputs):
```

```
        y = tf.matmul(inputs, self.W) + self.b
```

```
        if self.activation is not None:
```

```
            y = self.activation(y)
```

```
        return y
```

We can even define custom metrics and custom loss functions using the subclassing API. Refer to Section 7.3.1 of Chollet for details.

## ✓ Building the model using custom dense layer and functional API

```
inputs = Input(shape=(28*28,))
```

```
features = Custom_Dense(512, activation=tf.nn.relu)(inputs)
```

```
features = Custom_Dense(128, activation=tf.nn.relu)(features)
```

```
outputs = Custom_Dense(10, activation=tf.nn.softmax)(features)
```

```
model = keras.Model(inputs, outputs)
```

```
model.summary()
```

Model: "model\_3"

Layer (type)	Output Shape	Param #
=====		
input_3 (InputLayer)	[(None, 784)]	0
custom_dense (Custom_Dense )	(None, 512)	401920
custom_dense_1 (Custom_Den se)	(None, 128)	65664
custom_dense_2 (Custom_Den se)	(None, 10)	1290
=====		
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

```
plot_model(model)
```



```
model.compile(optimizer =keras.optimizers.RMSprop(),
              loss = keras.losses.SparseCategoricalCrossentropy(),
              metrics = ["accuracy"])
```

```
train_x = train_images[10000:]
train_y = train_labels[10000:]
val_x = train_images[:10000]
val_y = train_labels[:10000]
```

```
history = model.fit(x=train_x, y=train_y, epochs=10,
                    validation_data=(val_x, val_y))
```

```

Epoch 1/10
1563/1563 [=====] - 6s 4ms/step - loss: 0.2059 - accuracy: 0.9377 - val_loss: 0.1166 - val_accuracy: 0.9666
Epoch 2/10
1563/1563 [=====] - 6s 4ms/step - loss: 0.0939 - accuracy: 0.9733 - val_loss: 0.0929 - val_accuracy: 0.9733
Epoch 3/10
1563/1563 [=====] - 6s 4ms/step - loss: 0.0682 - accuracy: 0.9809 - val_loss: 0.0884 - val_accuracy: 0.9773
Epoch 4/10
1563/1563 [=====] - 6s 4ms/step - loss: 0.0569 - accuracy: 0.9854 - val_loss: 0.1114 - val_accuracy: 0.9751
Epoch 5/10
1563/1563 [=====] - 5s 3ms/step - loss: 0.0451 - accuracy: 0.9882 - val_loss: 0.1030 - val_accuracy: 0.9783
Epoch 6/10
1563/1563 [=====] - 6s 4ms/step - loss: 0.0381 - accuracy: 0.9902 - val_loss: 0.1444 - val_accuracy: 0.9751
Epoch 7/10
1563/1563 [=====] - 6s 4ms/step - loss: 0.0312 - accuracy: 0.9922 - val_loss: 0.1301 - val_accuracy: 0.9766
Epoch 8/10
1563/1563 [=====] - 6s 4ms/step - loss: 0.0291 - accuracy: 0.9929 - val_loss: 0.1377 - val_accuracy: 0.9796
Epoch 9/10
1563/1563 [=====] - 6s 4ms/step - loss: 0.0258 - accuracy: 0.9936 - val_loss: 0.1566 - val_accuracy: 0.9795
Epoch 10/10
1563/1563 [=====] - 5s 4ms/step - loss: 0.0222 - accuracy: 0.9949 - val_loss: 0.1730 - val_accuracy: 0.9775

```

```
data = pd.DataFrame(history.history)
data.head()
```

```


loss accuracy val_loss val_accuracy
0 0.205926 0.93768 0.116598 0.9660
1 0.093874 0.97334 0.092927 0.9733
2 0.068189 0.98088 0.088432 0.9773
3 0.056886 0.98538 0.111375 0.9751
4 0.045090 0.98816 0.102987 0.9783

```

```

plt.plot(range(1,11),data['loss'], label="Training Loss")
plt.plot(range(1,11),data['val_loss'],label="validation Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()

```



Q. What is the difference between evaluate() and predict()

A: evaluate() returns the loss score and evaluation score. predict() runs a forward pass for the given input data.

```
model.evaluate(test_images, test_labels)
```

#HW: LOAD the best saved model by using tf.load\_model and evaluate the loaded model

```
313/313 [=====] - 1s 2ms/step - loss: 0.1546 - accuracy: 0.9799
[0.15462972223758698, 0.9799000024795532]
```

```
class_predicted = np.argmax(model.predict(test_images), axis=1)
accuracy = np.sum(class_predicted == test_labels) / len(test_labels)
print(accuracy)
```

```
313/313 [=====] - 1s 2ms/step
0.9799
```

```
from sklearn.metrics import confusion_matrix, classification_report
```

```
print(classification_report(test_labels, class_predicted))
```

```
precision    recall  f1-score   support

0           0.97       0.99       0.98        980
1           0.99       0.99       0.99       1135
2           0.99       0.97       0.98       1032
3           0.97       0.99       0.98       1010
4           0.98       0.98       0.98        982
5           0.99       0.97       0.98        892
6           0.98       0.98       0.98        958
7           0.98       0.97       0.98       1028
8           0.97       0.97       0.97        974
9           0.97       0.98       0.97       1009

accuracy          0.98       10000
macro avg         0.98       0.98       0.98       10000
weighted avg      0.98       0.98       0.98       10000
```

```
print(confusion_matrix(test_labels, class_predicted))
```

```
[[ 973   1   0   1   1   0   2   1   1   0]
 [ 1125   1   1   0   0   2   1   4   0]
 [ 7   3  996   2   3   0   3   8  10   0]
 [ 2   1   1  996   0   0   0   1   3   6]
 [ 1   0   1   1  964   0   4   2   1   8]
 [ 3   0   0  12   2  868   4   0   2   1]
 [ 5   3   0   0   5   2  941   0   1   1]
 [ 2   3   3   3   2   1   0 1001   3  10]
 [ 6   1   2   5   1   6   0   1  949   3]
 [ 2   4   0   4   8   2   0   2   1  986]]
```

```
#save model
```

```
# https://www.tensorflow.org/api_docs/python/tf/keras/models/save_model
```

## ✓ Using Callbacks

A callback is a powerful tool to customize the behavior of a Keras model during training, evaluation, or inference.

```
# build model using functional API
inputs = Input(shape=(28*28,))
features = Dense(512, activation="relu")(inputs)
```

```
features = keras.layers.Dropout(0.5)(features)
outputs = Dense(10,activation="softmax")(features)
```

```
mnist_model = keras.Model(inputs, outputs)
```

```
from keras.callbacks import EarlyStopping, ModelCheckpoint, TensorBoard
```

```
callbacks_list = [EarlyStopping(monitor="val_loss", patience=2),
                  ModelCheckpoint("mnist_model_checkpoint", save_best_only=True),
                  TensorBoard(log_dir="/tensorboard_files")]
```

```
mnist_model.compile(optimizer =keras.optimizers.Adam(),
                    loss = keras.losses.SparseCategoricalCrossentropy(),
                    metrics = ["accuracy"])
```

```
mnist_model.fit(x=train_x, y=train_y, epochs=10,
                validation_data=(val_x, val_y),
                callbacks=callbacks_list,)
```

```
Epoch 1/10
1563/1563 [=====] - 6s 4ms/step - loss: 0.2884 - accuracy: 0.9147 - val_loss: 0.1436 - val_accuracy: 0.9575
Epoch 2/10
1563/1563 [=====] - 5s 4ms/step - loss: 0.1462 - accuracy: 0.9543 - val_loss: 0.1022 - val_accuracy: 0.9706
Epoch 3/10
1563/1563 [=====] - 5s 3ms/step - loss: 0.1126 - accuracy: 0.9644 - val_loss: 0.0913 - val_accuracy: 0.9728
Epoch 4/10
1563/1563 [=====] - 6s 4ms/step - loss: 0.0962 - accuracy: 0.9697 - val_loss: 0.0865 - val_accuracy: 0.9735
Epoch 5/10
1563/1563 [=====] - 6s 4ms/step - loss: 0.0839 - accuracy: 0.9737 - val_loss: 0.0801 - val_accuracy: 0.9754
Epoch 6/10
1563/1563 [=====] - 6s 4ms/step - loss: 0.0746 - accuracy: 0.9761 - val_loss: 0.0768 - val_accuracy: 0.9774
Epoch 7/10
1563/1563 [=====] - 8s 5ms/step - loss: 0.0678 - accuracy: 0.9783 - val_loss: 0.0770 - val_accuracy: 0.9765
Epoch 8/10
1563/1563 [=====] - 7s 4ms/step - loss: 0.0618 - accuracy: 0.9798 - val_loss: 0.0823 - val_accuracy: 0.9778
<keras.callbacks.History at 0x7fe4e3101760>
```

```
%load_ext tensorboard
```

```
%tensorboard --logdir /tensorboard_files
```





TensorBoard

SCALARS

GRAPHS

TIME SERIES

INACTIVE

## Reference

- Chollet, F. (2021). Deep learning with python. Manning Publications.

```

                                epoch accuracy
mnist_model.fit(x=train_x, y=train_y, epochs=10,
                # validation_data=(val_x, val_y),
                callbacks=callbacks_list,)

```



Epoch 1/10

```
1554/1563 [=====>.] - ETA: 0s - loss: 0.0573 - accuracy: 0.9810WARNING:tensorflow:Early stopping conditioned
WARNING:tensorflow:Can save best model only with val_loss available, skipping.
```

```
1563/1563 [=====] - 4s 3ms/step - loss: 0.0572 - accuracy: 0.9810
```

Epoch 2/10

```
1557/1563 [=====>.] - ETA: 0s - loss: 0.0545 - accuracy: 0.9821WARNING:tensorflow:Early stopping conditioned
WARNING:tensorflow:Can save best model only with val_loss available, skipping.
```

```
1563/1563 [=====] - 4s 3ms/step - loss: 0.0544 - accuracy: 0.9821
```

Epoch 3/10

```
1559/1563 [=====>.] - ETA: 0s - loss: 0.0499 - accuracy: 0.9837WARNING:tensorflow:Early stopping conditioned
WARNING:tensorflow:Can save best model only with val_loss available, skipping.
```

```
1563/1563 [=====] - 4s 3ms/step - loss: 0.0500 - accuracy: 0.9837
```

Epoch 4/10

```
1555/1563 [=====>.] - ETA: 0s - loss: 0.0500 - accuracy: 0.9832WARNING:tensorflow:Early stopping conditioned
WARNING:tensorflow:Can save best model only with val_loss available, skipping.
```

```
1563/1563 [=====] - 4s 3ms/step - loss: 0.0499 - accuracy: 0.9832
```

Epoch 5/10

```
1562/1563 [=====>.] - ETA: 0s - loss: 0.0462 - accuracy: 0.9849WARNING:tensorflow:Early stopping conditioned
WARNING:tensorflow:Can save best model only with val_loss available, skipping.
```

```
1563/1563 [=====] - 4s 3ms/step - loss: 0.0462 - accuracy: 0.9849
```

Epoch 6/10

```
191/1563 [==>.....] - ETA: 3s - loss: 0.0450 - accuracy: 0.9853
```

KeyboardInterrupt

Traceback (most recent call last)

```
<ipython-input-65-0c17a4dd562f> in <module>
```

```
----> 1 mnist_model.fit(x=train_x, y=train_y, epochs=10,
      2                 # validation_data=(val_x, val_y),
```